

Department of Electrical and Computer Engineering

University of Colorado at Boulder

ECEN5623 - Real Time Embedded Systems



Exercise 4

Submitted by

Parth T — Parth K

Submitted on March 25, 2024

Contents

List of Figures	1
List of Tables	1
1 Question 1	2
2 Question 2	3
3 Question 3	4
4 Question 4	8
5 Question 5	13
6 Question 6	19
7 References	20
Appendices	21
A C/Cpp/Python Code for the Implementation	21

List of Figures

1 lsusb	2
2 lsmod	2
3 Kernel Message	3
4 Cheese Version	3
5 Cheese Output	4
6 Build 1	5
7 Build 2	6
8 Build 3	6
9 Build 4	7
10 Canny	9
11 Canny Top	9
12 Hough Line	11
13 Hough Line Top	11
14 Hough Circle	12
15 Hough Circle top	13
16 Flowchart	14
17 Graph 1	18
18 Graph 2	18
19 Graph 3	19
20 Graph 4	19

List of Tables

*PDF is clickable

1 Question 1

Q: [10 points] Obtain a Logitech C200 or C270 camera or equivalent and verify that it is detected by the DE1-SoC, Raspberry Pi or Jetson Board USB driver. You can check the camera out from eStore or purchase one of your own, or use another camera that has a compliant UVC driver. Use lsusb, lsmod and dmesg kernel driver configuration tool to make sure your Logitech C2xx USB camera is plugged in and recognized by your DE1-SoC, Raspberry Pi or Jetson (note that on the Jetson, it does not use driver modules, but rather a monolithic kernel image, so lsmod will not look the same as other Linux systems – see what you can find by exploring /cat/proc on your Jetson to find the camera USB device). For the Jetson, do lsusb — grep C200 (or C270) and prove to the TA (and more importantly yourself) with that output (screenshot) that your camera is recognized. For systems other than a Jetson, do lsmod — grep video and verify that the UVC driver is loaded as well (<http://www.ideasonboard.org/uvc/>). To further verify, or debug if you don't see the UVC driver loaded in response to plugging in the USB camera, do dmesg — grep video or just dmesg and scroll through the log messages to see if your USB device was found. Capture all output and annotate what you see with descriptions to the best of your understanding.

Answer: Verify Camera Detection

- (a) Check USB Devices:

Command: lsusb

This command lists all USB devices currently connected to the system. Running this in the terminal helps you identify whether the Logitech camera is physically recognized by the system. The output includes details like the device ID and the manufacturer, which can be used to confirm the presence of the camera.

```
[...ta/university/RTEs_ECEN5623/Assignments/Excercise_4/answers/q4]
x parth ➤ lsusb | grep C270
Bus 003 Device 022: ID 046d:0825 Logitech, Inc. Webcam C270
```

Figure 1: lsusb

- (b) Verify UVC Driver Load:

Command: lsmod — grep video

Purpose: The lsmod command displays a list of loaded kernel modules, and piping this output to grep video filters the list to show only modules related to video, such as the UVC (USB Video Class) driver, often indicated by uvcvideo. This command verifies that the necessary driver for operating the camera is loaded in the system. The UVC driver is crucial because it provides generic support for USB video devices, ensuring that compliant cameras can work with a broad set of operating systems and applications without requiring proprietary drivers.

```
parth ➤ lsmod | grep uvcvideo
uvcvideo                  135168  0
videobuf2_vmalloc          20480   1 uvcvideo
uvc                      12288   1 uvcvideo
videobuf2_v4l2             40960   1 uvcvideo
videodev                  364544   2 videobuf2_v4l2, uvcvideo
videobuf2_common           86016   4 videobuf2_vmalloc, videobuf2_v4l2, uvcvideo, videobuf2_memops
mc                        86016   5 videodev, snd_usb_audio, videobuf2_v4l2, uvcvideo, videobuf2_common
```

Figure 2: lsmod

(c) Inspect Kernel Messages:

Commands: dmesg — grep video and dmesg

Purpose: These commands are used to inspect kernel messages. dmesg displays a log of system messages, including hardware detection, drivers loaded, and errors encountered during the boot process or while the system is running. dmesg — grep video narrows down this log to entries related to video devices and drivers. This can reveal whether the camera was successfully initialized, recognized, and configured by the kernel, or if there were any issues during its detection process.

```
[..ta/university/RTES_ECEN5623/Assignments/Excercise_4/answers/q4]
parth ➤ sudo dmesg | tail -f
[31994.011447] r8152 5-2.3:1.0 enx00e04c3606f4: renamed from eth0
[31994.100320] kauditd_printk_skb: 14 callbacks suppressed
[31994.100323] audit: type=1107 audit(1711329632.071:1166): pid=852 uid=102 auid=4294967295 ses=4294967295 subj=unconfined msg='apparmor="DENIED" operation="dbus_signal" bus="system" path="/org/freedesktop/NetworkManager/Settings" interface="org.freedesktop.NetworkManager.Settings" member="NewConnection" name=:1.3 mask="receive" pid=28429 label="snap.zoom-client.zoom-client" peer_pid=857 peer_label="unconfined"
          exe="/usr/bin/dbus-daemon" sauid=102 hostname=? addr=? terminal=?
[32300.540322] usb 3-2: USB disconnect, device number 21
[32303.416323] usb 3-2: new high-speed USB device number 22 using xhci_hcd
[32303.784761] usb 3-2: New USB device found, idVendor=046d, idProduct=0825, bcdDevice= 0.12
[32303.784777] usb 3-2: New USB device strings: Mfr=0, Product=0, SerialNumber=2
[32303.784784] usb 3-2: SerialNumber: AB20F2B0
[32303.848058] usb 3-2: Found UVC 1.00 device <unnamed> (046d:0825)
```

Figure 3: Kernel Message

2 Question 2

Q: Option 2: Cheese If you do not have cheese, do sudo apt-get install cheese on your Jetson, Raspberry Pi or DE1-SoC board or other native Linux system. This should not only install nice camera capture GUI tools, but also the V4L2 API described well in this series of Linux articles - <http://lwn.net/Articles/203924/> . Running cheese should provide an interactive camera control session for your Logitech C2xx camera – if you have issues connecting to your camera do a “man cheese” and specify your camera device file entry point (e.g. /dev/video0). Show that you tested your camera with a cheese screen dump and test photo.

Answer: Install Cheese: Execute the following command:

```
1 sudo apt-get install cheese
2
```

```
[..ta/university/RTES_ECEN5623/Assignments/Excercise_4/answers/q4]
parth ➤ cheese --version
Cheese 41.1
```

Figure 4: Cheese Version

Manual Device Selection: If Cheese does not automatically detect your camera, or if you encounter issues, you can specify the camera device file manually:

Determine your camera’s device file (/dev/video0 for the first detected video device). You can list video devices with ls /dev/video*. Launch Cheese with a specified device file by running: cheese –device=/dev/video0 (replace /dev/video0 with camera’s device file if different).

Screen Dump: screenshot of the Cheese application window showing the live camera feed with keyboard shortcut (e.g., PrtScn key).



Figure 5: Cheese Output

3 Question 3

Q: [10 points] Using your verified Logitech C2xx camera on a DE1-SoC, Raspberry Pi or Jetson, verify that it can stream continuously to a raw image buffer for transformation and processing using any one of the below methods:

Answer:

- (a) Update Package Index and Install Prerequisites: This step involves updating the package index on your Linux system and installing the minimal prerequisites required to compile OpenCV, such as cmake, g++, wget, and unzip.
- (b) Download OpenCV Sources: The OpenCV source code is downloaded from its GitHub repository. The wget command is used to fetch the zip archive of the OpenCV source code designated by the 4.x version tag, which is then unpacked using unzip.
- (c) Create a Build Directory: A separate build directory is created to keep the compilation process organized and separate from the source code.
- (d) Configure the Build: The cmake command is used to configure the build environment. This step prepares the makefile based on the available source code and specified options.
- (e) Compile OpenCV: The compilation process is initiated with the cmake -build . command. This step compiles the source code into executable binaries and libraries.

```

1 # Install minimal prerequisites (Ubuntu 18.04 as reference)
2 sudo apt update && sudo apt install -y cmake g++ wget unzip
3
4 # Download and unpack sources
5 wget -O opencv.zip https://github.com/opencv/opencv/archive/4.x.zip
6 unzip opencv.zip
7
8 # Create build directory
9 mkdir -p build && cd build
10
11 # Configure
12 cmake .. / opencv-4.x
13
14 # Build
15 cmake --build .
16
17 # make
18 make -j4
19
20 # make install
21 sudo make install
22

```

The screenshot shows a terminal window titled "parth@rog:~/build". The window displays the output of a "make install" command. The output lists numerous OpenCV targets being built, each accompanied by a progress bar indicating its completion percentage. The targets include libwebp, libjasper, ippw, libprotobuf, quirc, ittnotify, ade, opencv_videoio_plugins, opencv_core, opencv_imgproc, opencv_imgcodecs, opencv_videoio, opencv_highgui, opencv_ts, opencv_test_core, opencv_perf_core, opencv_flann, opencv_test_flann, opencv_test_imgproc, opencv_perf_imgproc, opencv_ml, opencv_test_ml, opencv_photo, opencv_test_photo, opencv_perf_photo, opencv_dnn, opencv_test_dnn, opencv_perf_dnn, opencv_features2d, opencv_test_features2d, opencv_perf_features2d, opencv_gapi, opencv_test_gapi, opencv_perf_gapi, opencv_test_imgcodecs, opencv_perf_imgcodecs, opencv_perf_imgcodecs, opencv_test_videoio, and opencv_gapi. The build process is nearly complete, reaching approximately 81% completion.

```

[~/build]
parth > sudo make install
[sudo] password for parth:
 9% Built target libwebp
12% Built target libjasper
14% Built target ippw
20% Built target libprotobuf
21% Built target quirc
21% Built target ittnotify
22% Built target ade
22% Built target opencv_videoio_plugins
30% Built target opencv_core
36% Built target opencv_imgproc
38% Built target opencv_imgcodecs
39% Built target opencv_videoio
39% Built target opencv_highgui
40% Built target opencv_ts
43% Built target opencv_test_core
46% Built target opencv_perf_core
46% Built target opencv_flann
46% Built target opencv_test_flann
50% Built target opencv_test_imgproc
53% Built target opencv_perf_imgproc
54% Built target opencv_ml
54% Built target opencv_test_ml
55% Built target opencv_photo
56% Built target opencv_test_photo
57% Built target opencv_perf_photo
65% Built target opencv_dnn
67% Built target opencv_test_dnn
67% Built target opencv_perf_dnn
69% Built target opencv_features2d
70% Built target opencv_test_features2d
71% Built target opencv_perf_features2d
75% Built target opencv_gapi
Consolidate compiler generated dependencies of target opencv_test_gapi
[ 79%] Built target opencv_test_gapi
[ 80%] Built target opencv_perf_gapi
Consolidate compiler generated dependencies of target opencv_test_imgcodecs
[ 80%] Built target opencv_test_imgcodecs
Consolidate compiler generated dependencies of target opencv_perf_imgcodecs
[ 80%] Built target opencv_perf_imgcodecs
Consolidate compiler generated dependencies of target opencv_test_videoio
[ 81%] Built target opencv_test_videoio

```

Figure 6: Build 1

```

-- Up-to-date: /usr/local/lib/libopencv_video.so.4.1
-- Up-to-date: /usr/local/lib/libopencv_video.so
-- Up-to-date: /usr/local/include/opencv4/opencv2/video/background_segm.hpp
-- Up-to-date: /usr/local/include/opencv4/opencv2/video/legacy/constants_c.h
-- Up-to-date: /usr/local/include/opencv4/opencv2/video/tracking.hpp
-- Up-to-date: /usr/local/include/opencv4/opencv2/video/video.hpp
-- Up-to-date: /usr/local/share/java/opencv4/libopencv_java411.so
-- Up-to-date: /usr/local/share/java/opencv4/opencv-411.jar
-- Up-to-date: /usr/local/lib/python3.10/dist-packages/cv2/_init__.py
-- Up-to-date: /usr/local/lib/python3.10/dist-packages/cv2/load_config_py2.py
-- Up-to-date: /usr/local/lib/python3.10/dist-packages/cv2/load_config_py3.py
-- Up-to-date: /usr/local/lib/python3.10/dist-packages/cv2/config.py
-- Up-to-date: /usr/local/lib/python3.10/dist-packages/cv2/python-3.10cv2.cpython-310-x86_64-linux-gnu.so
-- Up-to-date: /usr/local/lib/python3.10/dist-packages/cv2/config-3.10.py
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_eye.xml
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_eye_tree_eyeglasses.xml
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_frontalcatface.xml
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_frontalcatface_extended.xml
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_frontalface.alt.xml
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_frontalface.alt2.xml
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_frontalface.alt_tree.xml
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_frontalface_default.xml
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_fullbody.xml
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_lefteye_2splits.xml
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_licence_plate_rus_16stages.xml
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_lowerbody.xml
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_profileface.xml
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_righteye_2splits.xml
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_russian_plate_number.xml
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_smile.xml
-- Up-to-date: /usr/local/share/opencv4/haarcascades/haarcascade_upperbody.xml
-- Up-to-date: /usr/local/share/opencv4/lbpcascades/lbpcascade_frontalcatface.xml
-- Up-to-date: /usr/local/share/opencv4/lbpcascades/lbpcascade_frontalface.xml
-- Up-to-date: /usr/local/share/opencv4/lbpcascades/lbpcascade_frontalface_improved.xml
-- Up-to-date: /usr/local/share/opencv4/lbpcascades/lbpcascade_profileface.xml
-- Up-to-date: /usr/local/share/opencv4/lbpcascades/lbpcascade_silverware.xml
-- Up-to-date: /usr/local/bin/opencv_annotation
-- Up-to-date: /usr/local/bin/opencv_visualisation
-- Up-to-date: /usr/local/bin/opencv_interactive-calibration
-- Up-to-date: /usr/local/bin/opencv_version

[~/build] parth

```

Figure 7: Build 2

```

[~/build] parth > sudo make install
[sudo] password for parth:
  9% Built target libwebp
 12% Built target libjasper
14% Built target ippwi
20% Built target libprotobuf
21% Built target quirke
21% Built target ittnotify
22% Built target ade
22% Built target opencv_videoio_plugins
30% Built target opencv_core
36% Built target opencv_imgproc
38% Built target opencv_imgcodecs
39% Built target opencv_videoio
39% Built target opencv_highgui
40% Built target opencv_ts
43% Built target opencv_test_core
46% Built target opencv_perf_core
46% Built target opencv_flann
46% Built target opencv_test_flann
50% Built target opencv_test_imgproc
53% Built target opencv_perf_imgproc
54% Built target opencv_ml
54% Built target opencv_test_ml
55% Built target opencv_photo
56% Built target opencv_test_photo
57% Built target opencv_perf_photo
65% Built target opencv_dnn
67% Built target opencv_test_dnn
67% Built target opencv_perf_dnn
69% Built target opencv_features2d
70% Built target opencv_test_features2d
71% Built target opencv_perf_features2d
75% Built target opencv_gapi
Consolidate compiler generated dependencies of target opencv_test_gapi
  79% Built target opencv_test_gapi
  80% Built target opencv_perf_gapi
Consolidate compiler generated dependencies of target opencv_test_imgcodecs
  80% Built target opencv_test_imgcodecs
Consolidate compiler generated dependencies of target opencv_perf_imgcodecs
  80% Built target opencv_perf_imgcodecs
Consolidate compiler generated dependencies of target opencv_test_videoio
  81% Built target opencv_test_videoio

```

Figure 8: Build 3

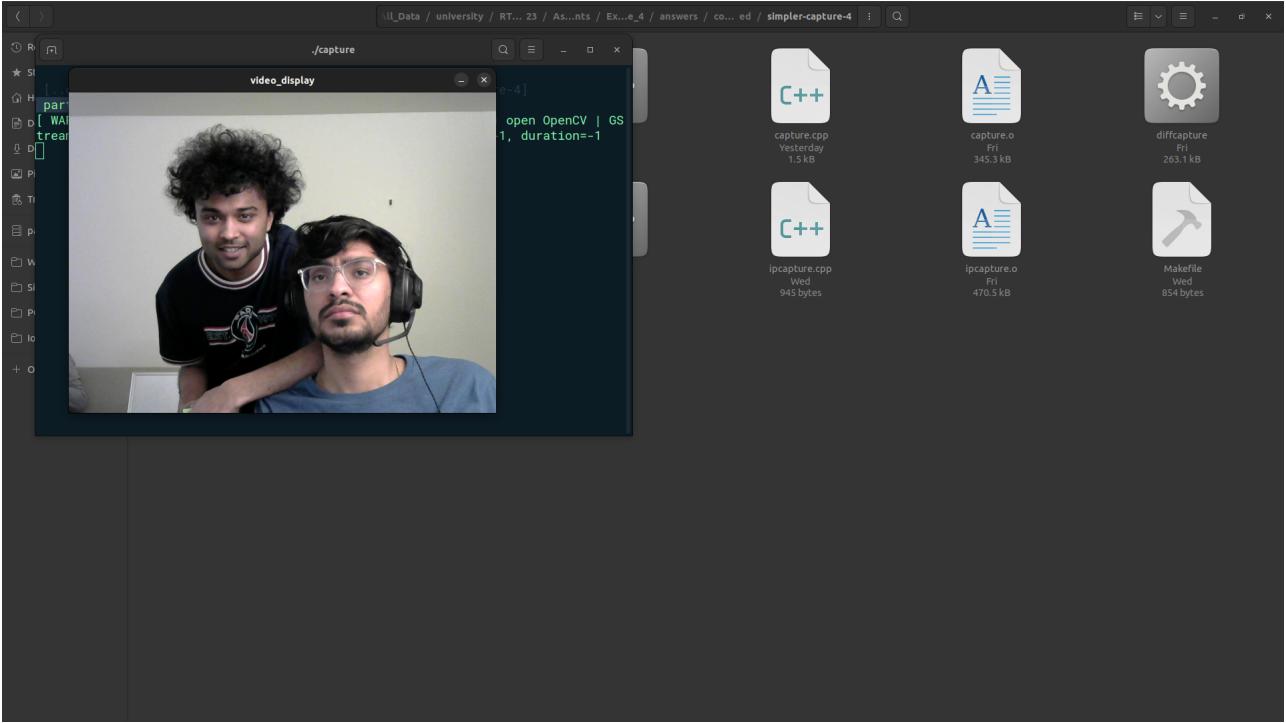


Figure 9: Build 4

- (a) `VideoCapture cam0(0);`: Creates an object `cam0` for capturing video from the camera. The parameter `(0)` specifies the first camera device. This class provides a way to interact with the video capturing functionality of OpenCV.
- (b) `namedWindow "video_display";`: Creates a window named `"video_display"` for displaying the video. This function is part of the High-level GUI module.
- (c) `cam0.isOpened();`: Checks if the video capture device has been successfully initialized.
- (d) `exit(SYSTEM_ERROR);`: Exits the program with a system-defined error status if the camera cannot be opened.
- (e) `cam0.setCAP_PROP_FRAME_WIDTH, 640;` and `cam0.setCAP_PROP_FRAME_HEIGHT, 480;`: Sets the properties of the capture device, specifically the frame width and height to `640x480` pixels. `CAP_PROP_FRAME_WIDTH` and `CAP_PROP_FRAME_HEIGHT` are predefined constants that specify which property to set.
- (f) `Mat frame640, 480, ::`: Defines a matrix (`Mat`) object named `frame` to store individual frames captured from the camera. The constructor parameters typically include the dimensions and the type of the matrix, but here it seems to be a typo or an incomplete line of code, as the type is missing.
- (g) `cam0.read(frame);`: Captures the next frame from the video capture device and stores it in `frame`. This method returns false if no frames have been captured.
- (h) `imshow "video_display", frame;`: Displays the frame in the window named `"video_display"`. This function is called within the loop to continuously update the window with new frames from the camera.
- (i) `waitKey(10);`: Waits for a key event for a brief period (10 milliseconds). This function is essential for the GUI to process events and update the window. The return value is stored in `winInput`.
- (j) `if ((winInput = waitKey(10)) == ESCAPE_KEY):`: Checks if the ESC key `ASCII value 27` has been pressed. If so, it breaks out of the loop, ending the video capture.
- (k) `destroyWindow "video_display";`: Closes the window named `"video_display"` before the program exits.

4 Question 4

Q: [15 points] Review POSIX-examples and especially **POSIX_MQ_LOOP** and build the code related to **POSIX message queues** and run them to learn about basic use.

Answer:

(a)

Canny Here's an explanation of the provided Code

- Mat variables canny_frame, timg_gray, timg_grad, and frame are declared to store image data.
- Integer variables lowThreshold, max_lowThreshold, ratio, and kernel_size are declared for Canny edge detection parameters.

CannyThreshold Function:

- This function is used to perform Canny edge detection on the captured frame.
- It converts the frame to grayscale using cvtColor().
- Noise reduction is applied using blur().
- Canny edge detection is performed using Canny().
- The result is displayed in the window_name window using imshow().

Main Function:

- Video capture object cam0 is initialized for the default camera (index 0).
- A window named "video_display" is created for displaying the captured video frames.
- The frame size is set to 640x480 pixels using cam0.set().
- Within a while loop:
 - Frames are continuously captured from the camera using cam0.read().
 - The captured frame is displayed in the "video_display" window using imshow().
 - The Canny edge detection trackbar is created and CannyThreshold function is called to apply edge detection on the frame.
 - The program exits the loop if the escape key (ESC) is pressed.

This code captures video frames from the camera, performs Canny edge detection, and

Cannycam

```

Activities CannyCam Mar 24 7:13 PM
Open ... Save ...
cannycam.cpp -/Work/All_Data/university/RTE...623/Assignments/Exercise_4/answers/q4
33 // Canny detector
34 Canny( canny_frame, canny_frame, lowThreshold, lowThreshold*ratio, kernel_size );
35
36 /// Using Canny's output as a mask, we display our result
37 timg_grad = Scalar::all(0);
38
39 frame.copyTo( timg_grad, canny_frame);
40
41 imshow( window_name, timg_grad );
42
43 }
44
45
46
47 int main( int a
48 {
49     VideoCapture
50     namedWindow(
51         char wInpu
52
53     if (cam0.is
54     {
55         ext(SYS
56     }
57
58     cam0.set(CAP
59     cam0.set(CAP
60
61     while (1)
62     {
63         cam0.read
64
65         imshow("v
66
67         namedWind
68         //create
69
70         CannyThre
71
72         if ((winInput == waitKey(0) == ESCAPE_KEY)
73             //if ((winInput == waitKey(0) == ESCAPE_KEY)
74             {
75                 break;
76             }
77             else if (winInput == 'n')
78             {
79                 printf("Input %c is ignored\n", winInput);
80             }
81
82     }
83
84     destroyWindow("video_display");
85
86 }

```

[WARN:0] global ./modules/videoio/src/cap_gstreamer.cpp (1100) open OpenCV | GStreamer warning: Cannot query video position: status=0, value=-1, duration=-1

C++ Tab Width: 8 Ln 59, Col 41 INS

Figure 10: Canny

Cannycam top

```

./cannycam
top - 23:18:29 up 3:22, 1 user, load average: 3.44, 4.01, 3.69
Tasks: 431 total, 4 running, 427 sleeping, 0 stopped, 0 zombie
%Cpu(s): 21.4 us, 1.6 sy, 0.0 ni, 75.5 id, 0.0 wa, 0.0 hi, 1.5 si, 0.0 st
MiB Mem : 23435.3 total, 12437.3 free, 6302.5 used, 4695.4 buff/cache
MiB Swap: 2048.0 total, 0.0 free, 0.0 used. 16420.2 avail Mem

PID USER PR NI VIRT RES SHR S %CPU %MEM TIME+ COMMAND
5143 partth 20 0 132544 54748 4344 R 100.0 0.2 201:09.84 pdflatex
15287 partth 20 0 134028 59012 4608 R 100.0 0.2 107:07.35 pdflatex
24045 partth 20 0 1751288 08756 65328 S 66.9 0.4 0:13.99 cannycam
2769 partth 20 0 4712864 546861 149308 R 18.9 2.3 9:59.12 ghome-s+
2458 partth 20 0 25.1g 253692 154404 S 8.6 1.1 9:03.77 Xorg
9224 partth 20 0 1131.4g 284896 194444 S 4.3 1.2 8:58.29 opera
967 root -51 0 0 0 0 S 3.3 0.0 4:41.42 irq/99+-+
5354 partth 20 0 32.8g 646416 204572 S 1.7 2.7 9:00.20 opera
18262 root 20 0 0 0 0 D 1.0 0.0 0:05.09 kworker+
23194 root 0 -20 0 0 0 I 1.0 0.0 0:00.28 kworker+
24107 partth 20 0 22000 4352 3328 R 1.0 0.0 0:00.11 top
1277 root -2 0 0 0 0 S 0.7 0.0 0:25.25 gfx_low
2944 partth 20 0 324136 11800 6912 S 0.7 0.0 0:16.09 ibus-da+
22490 root 20 0 0 0 0 D 0.7 0.0 0:04.04 kworker+
2911 partth 20 0 1266456 129292 67428 S 0.3 0.5 0:56.00 nautilus
3235 partth 20 0 172264 7296 6656 S 0.3 0.0 0:03.94 ibus-en+
3657 partth 20 0 32.7g 141016 108524 S 0.3 0.6 0:39.28 slack

```

Figure 11: Canny Top

Here's an explanation of the provided code

- Mat variables dst, cdst, and cdstP are declared to store images. default_file and filename are declared to store the default image file path and the actual image file path, respectively. src is declared to store the loaded image.

- Image Loading: `imread()` function is used to load an image from the specified file path (filename) in grayscale mode (`IMREAD_GRAYSCALE`).
- Edge Detection (Canny): `Canny()` function is applied to detect edges in the loaded image (src). The detected edges are stored in the dst matrix.
- Color Conversion: `cvtColor()` function converts the grayscale edge image (dst) to a BGR (Blue-Green-Red) color image (cdst) for visualization.
- Standard Hough Line Transform: `HoughLines()` function detects lines in the edge image (dst) using the Standard Hough
- Drawing Detected Lines: For each detected line, `line()` function draws a line on the output image (cdst) using the line parameters (rho and theta). Probabilistic Hough Line Transform: `HoughLinesP()` function detects lines in the edge image (dst) using the Probabilistic Hough Line Transform. Detected lines are stored in the linesP vector.
- Drawing Probabilistic Detected Lines: For each detected line, `line()` function draws a line on the output image (cdstP) using the line parameters (x1, y1, x2, y2). Additionally, lines are drawn on the original image (src) to visualize the detected lines.
- Displaying Results: `imshow()` function displays the original image (src) and the image with detected lines (cdstP) in separate windows. This code loads an image, performs edge detection using the Canny algorithm, detects lines in the edge image using both the Standard Hough Line Transform and the Probabilistic Hough Line Transform, and visualizes the detected lines on the original image. Finally, it displays the results in separate windows and waits for a key press before exiting.

Hough Line Transform

The following changes have been made to the `houghline.cpp` code to make it work with continuous video input in `houghlinecam.cpp`:

- Video Capture: The `houghlinecam.cpp` code includes the `opencv2/videoio.hpp` header file to access the video capture functionality. It creates a `VideoCapture` object `cam0` to read frames from the default camera (index 0).
- Continuous Video Processing Loop: The code enters a continuous loop where it reads frames from the camera using `cam0.read(frame)`, displays the frame in the "video_display" window using `imshow("video_display", frame)`, performs edge detection using the Canny algorithm, copies the edges to the output images (cdst and cdstP), and calls the `HoughLinePTransform()` function to detect and draw lines.
- Keyboard Input Handling: The loop checks for keyboard input using `waitKey(10)`, which waits for 10 milliseconds for a key press. If the ESC key is pressed, the loop breaks and the program exits. If the 'n' key is pressed, it prints a message indicating that the input is ignored.
- Hough Line Probabilistic Transform Function: The `HoughLinePTransform()` function is defined to encapsulate the Probabilistic Hough Line Transform operation. This function performs the transform on the dst image (containing the edges), draws the detected lines on the cdstP and src (original frame) images using the `line` function, and displays the result in a window named "Prob HL Transform".
- Window Destruction: After the continuous loop exits, the "video_display" window is closed using the `destroyWindow` function.

```

Activities Houghlinecam Mar 24 7:14 PM
Open / cannycam.cpp -/Work>All_Data/university/RTEs_ECE15623/Assignments/Exercise_4/answers/q4 Save
33 // Canny detector
34 Canny( canny_frame, canny_frame, lowThreshold, lowThreshold*ratio, kernel_size );
35
36 // Using Canny's output as a mask, we display our result
37 timg_grad = scalar::all(0);
38
39 frame.copyTo( timg_grad, canny_frame );
40 imshow( window_name, timg_grad );
41
42 }
43
44 int main( int argc, char** argv )
45 {
46 VideoCapture cam0();
47 namedWindow("video_display");
48 char winInput;
49
50 if ( !cam0.isOpened() )
51 {
52     exit( SYSTEM_ERROR );
53 }
54 cam0.set(CAP_PROP_FRAME_WIDTH, 640);
55 cam0.set(CAP_PROP_FRAME_HEIGHT, 480);
56
57 while ( 1 )
58 {
59     cam0.read(frame);
60     imshow("video_display", frame);
61     namedWindow( window_name, WINDOW_AUTOSIZE );
62     //createTrackbar( "Min Threshold:", window_name, &lowThreshold, max_lowThreshold );
63     CannyThreshold(0, 0);
64     if ((winInput = waitKey(1)) == ESCAPE_KEY)
65     {
66         break;
67     }
68     else if (winInput == 'n')
69     {
70         printf("Input %c is ignored\n", winInput);
71     }
72 }
73 destroyWindow("video_display");
74
75 }
76
77
78
79
80
81
82
83
84
85
86

```

C++ Tab Width: 8 Ln 59, Col 41 INS

Figure 12: Hough Line

Hough Line Transform top

```

Activities Terminal Mar 24 7:17 PM
Open / cannycam.cpp -/Work>All_Data/university/RTEs_ECE15623/Assignments/Exercise_4/answers/q4 Save
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86

```

```

top - 19:17:24 up 16:21, 1 user, load average: 2.32, 2.17, 1.73
Tasks: 466 total, 2 running, 464 sleeping, 0 stopped, 0 zombie
%Cpu(s): 19.9 us, 2.7 sy, 0.0 ni, 76.9 id, 0.0 wa, 0.0 hi, 0.6 si, 0.0 st
MiB Mem : 23435.3 total, 7383.2 free, 9512.6 used, 6539.5 buff/cache
MiB Swap : 2048.0 total, 0.0 free, 0.0 used. 13233.8 avail Mem

PID USER PR NI VIRT RES SHR S %CPU %MEM TIME+ COMMAND
57634 parth 20 0 1754552 107044 65772 R 69.3 0.4 0:11.31 houghlinecam
7018 parth 20 0 1131.5g 274564 182876 S 61.1 1.1 286:54.35 opera
22465 parth 20 0 32.8g 492708 203412 S 28.1 2.1 92:35.44 opera
43207 parth 20 0 1145.9g 299028 179444 S 13.5 1.2 5:10.53 opera
2245 parth 20 0 25.1g 298624 128576 S 12.5 1.2 24:31.28 xorg
2559 parth 20 0 5149604 756404 154776 S 11.9 3.2 23:12.78 gnome-shell
853 root -51 0 0 0 0 I 5.9 0.0 28:12.37 irq/99-nvidia
2775 parth 20 0 642544 29952 21772 S 2.0 0.1 0:05.39 gsd-color
28429 parth 20 0 6826832 993252 364668 S 1.7 4.1 37:32.16 zoom
55137 root 0 -26 0 0 0 I 1.7 0.0 0:01.44 kworker/u33:1-uvcvideo
57576 parth 20 0 22136 4352 3328 R 1.0 0.0 0:00.39 top
2201 parth 20 0 10264 6392 4184 S 0.7 0.0 0:08.21 dbus-daemon
3640 parth 20 0 1131.5g 322824 112852 S 0.7 1.3 3:43.59 slack
55909 root 20 0 0 0 0 I 0.7 0.0 0:00.82 kworker/u32:0-writeback
56286 root 20 0 0 0 0 I 0.7 0.0 0:00.82 kworker/u32:4-events_un-
16 root 20 0 0 0 0 I 0.3 0.0 0:42.35 rcu_preempt
358 root 20 0 0 0 0 S 0.3 0.0 0:03.87 jbd2/vmme0n1p2-8
852 message+ 20 0 11352 7192 4212 S 0.3 0.0 0:09.38 dbus-daemon
857 root 20 0 270320 20000 16000 S 0.3 0.1 0:23.26 NetworkManager
1226 root -2 0 0 0 0 S 0.3 0.0 0:53.29 gfx_low
2862 colord 20 0 254292 13484 9660 S 0.3 0.1 0:01.49 colord

```

C++ Tab Width: 8 Ln 59, Col 41 INS

Figure 13: Hough Line Top

Hough Circle Transform

Here's an explanation of the provided code

- Image Loading: This code (main() function) loads an image specified by the user or uses a default image ("..../data/smarties.png") if no argument is provided.
- Image Processing:

- The loaded image (src) is directly converted to grayscale using cvtColor() and then median blur is applied to the grayscale image.
- Circle Detection: Detect circles in the processed grayscale image using the Hough Circle Transform (HoughCircles()) function.
- Circle Drawing: After detecting circles, draw circles on the original color image (src) using the circle() function. Image Display: Finally, display the image with detected circles using imshow() and wait for a key press using waitKey() before exiting.

The screenshot shows a terminal window titled 'Activities Houghcirccam' containing C++ code for Hough Circle detection. The code includes imports for cv, cvtColor, medianBlur, and HoughCircles. It initializes a VideoCapture object, sets its properties, and enters a loop where it reads frames, converts them to grayscale, applies median blur, and uses HoughCircles to detect circles. It then draws these circles on the original frame and displays it in a window named 'video_display'. The terminal also shows command-line arguments and some warning messages from OpenCV.

```

Activities Houghcirccam
Mar 24 7:14 PM
cannycam.cpp
-/Work/All_Data/university/RTEs_ECE15623/Assignments/Exercise_4/answers/q4

33 // Canny detector
34 Canny( canny_frame, canny_frame, lowThreshold, lowThreshold*ratio, kernel_size );
35
36 // Using Canny's output as a mask, we display our result
37 timg_grad = Scalar::all(0);
38
39 frame.copyTo( timg_grad, canny_frame );
40
41 imshow( window_name, timg_grad );
42
43 }
44
45
46
47 int main( int argc,
48 {
49 VideoCapture cam0;
50 namedWindow("video_display");
51 char winInput;
52
53 if ( !cam0.isOpened() )
54 {
55 exit( SYSTEM_ERROR );
56 }
57
58 cam0.set(CAP_PROP_FRAME_WIDTH, 640);
59 cam0.set(CAP_PROP_FRAME_HEIGHT, 480);
60
61 while (1)
62 {
63 cam0.read(frame);
64 imshow("video_display", frame);
65 namedWindow("detected circles");
66 //createTrackbar("CannyThresh", "detected circles", &can
67 CannyThreshold, 255);
68
69 if ((winInput = waitKey(10)) == ESCAPE_KEY)
70 //if ((winInput = waitKey(0)) == ESCAPE_KEY)
71 {
72 break;
73 }
74 else if (winInput == 'n')
75 {
76 printf("Input %c is ignored\n", winInput);
77 }
78
79
80
81 }
82
83 destroyWindow("video_display");
84
85
86 }

C++ Tab Width: 8 Ln 59, Col 41 INS

```

Figure 14: Hough Circle

Hough Circle Transform top

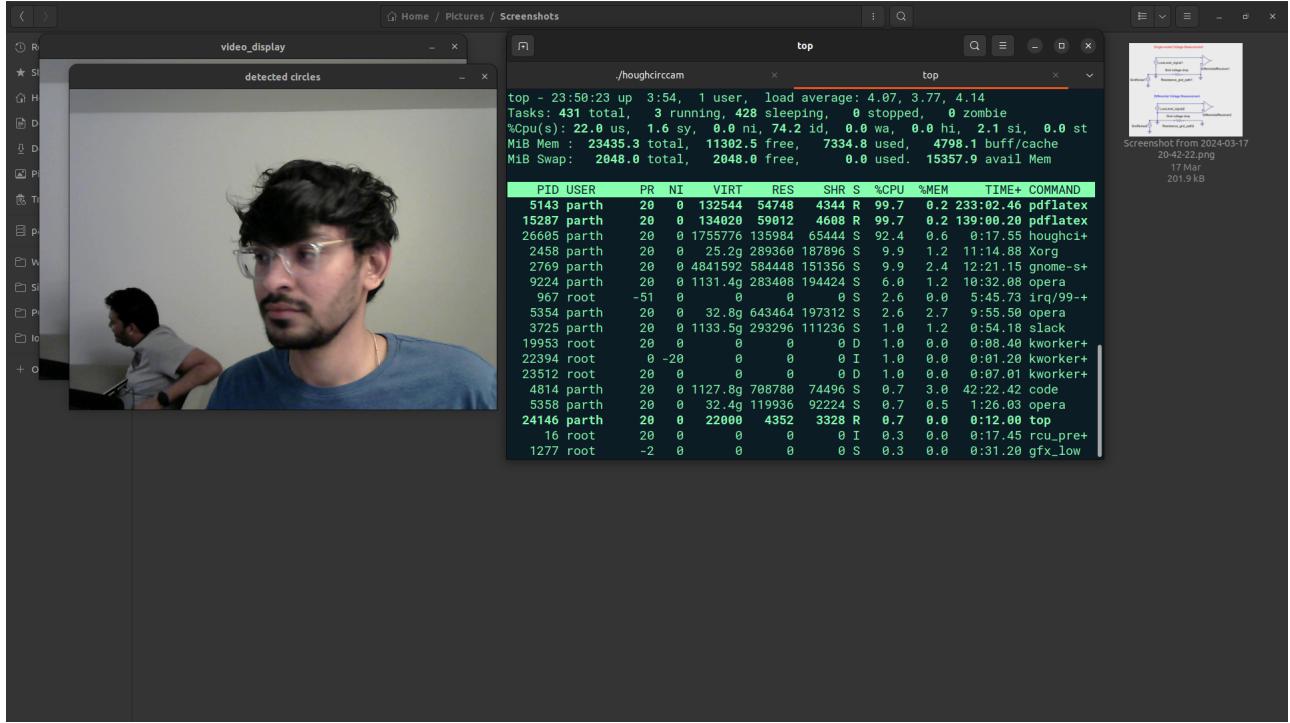


Figure 15: Hough Circle top

5 Question 5

Q: Using a Logitech C2xx, choose 3 real-time interactive transformations to compare in terms of average frame rate at a given resolution for a range of at least 3 resolutions in a common aspect ratio (e.g. 4:3 for 1280x960, 640x480, 320x240, 160x120, 80x60) by adding time-stamps (there should be a logging thread) to analyze the potential throughput. You should then get at least 9 separate datasets running 1 transformation at a time. Based on average analysis, pick a reasonable soft real-time deadline e.g. if average framerate is 12Hz, choose a deadline of 100 milliseconds to provide some margin and convert the processing to SCHED_FIFO and determine if you can meet deadlines with predictability and measure statistical jitter in the frame rate relative to your deadline for both the default scheduler and SCHED_FIFO in each case.

Answer:

(a) Design Concepts (Flow Chart or Diagram)

- Flow Chart/Diagram Creation:

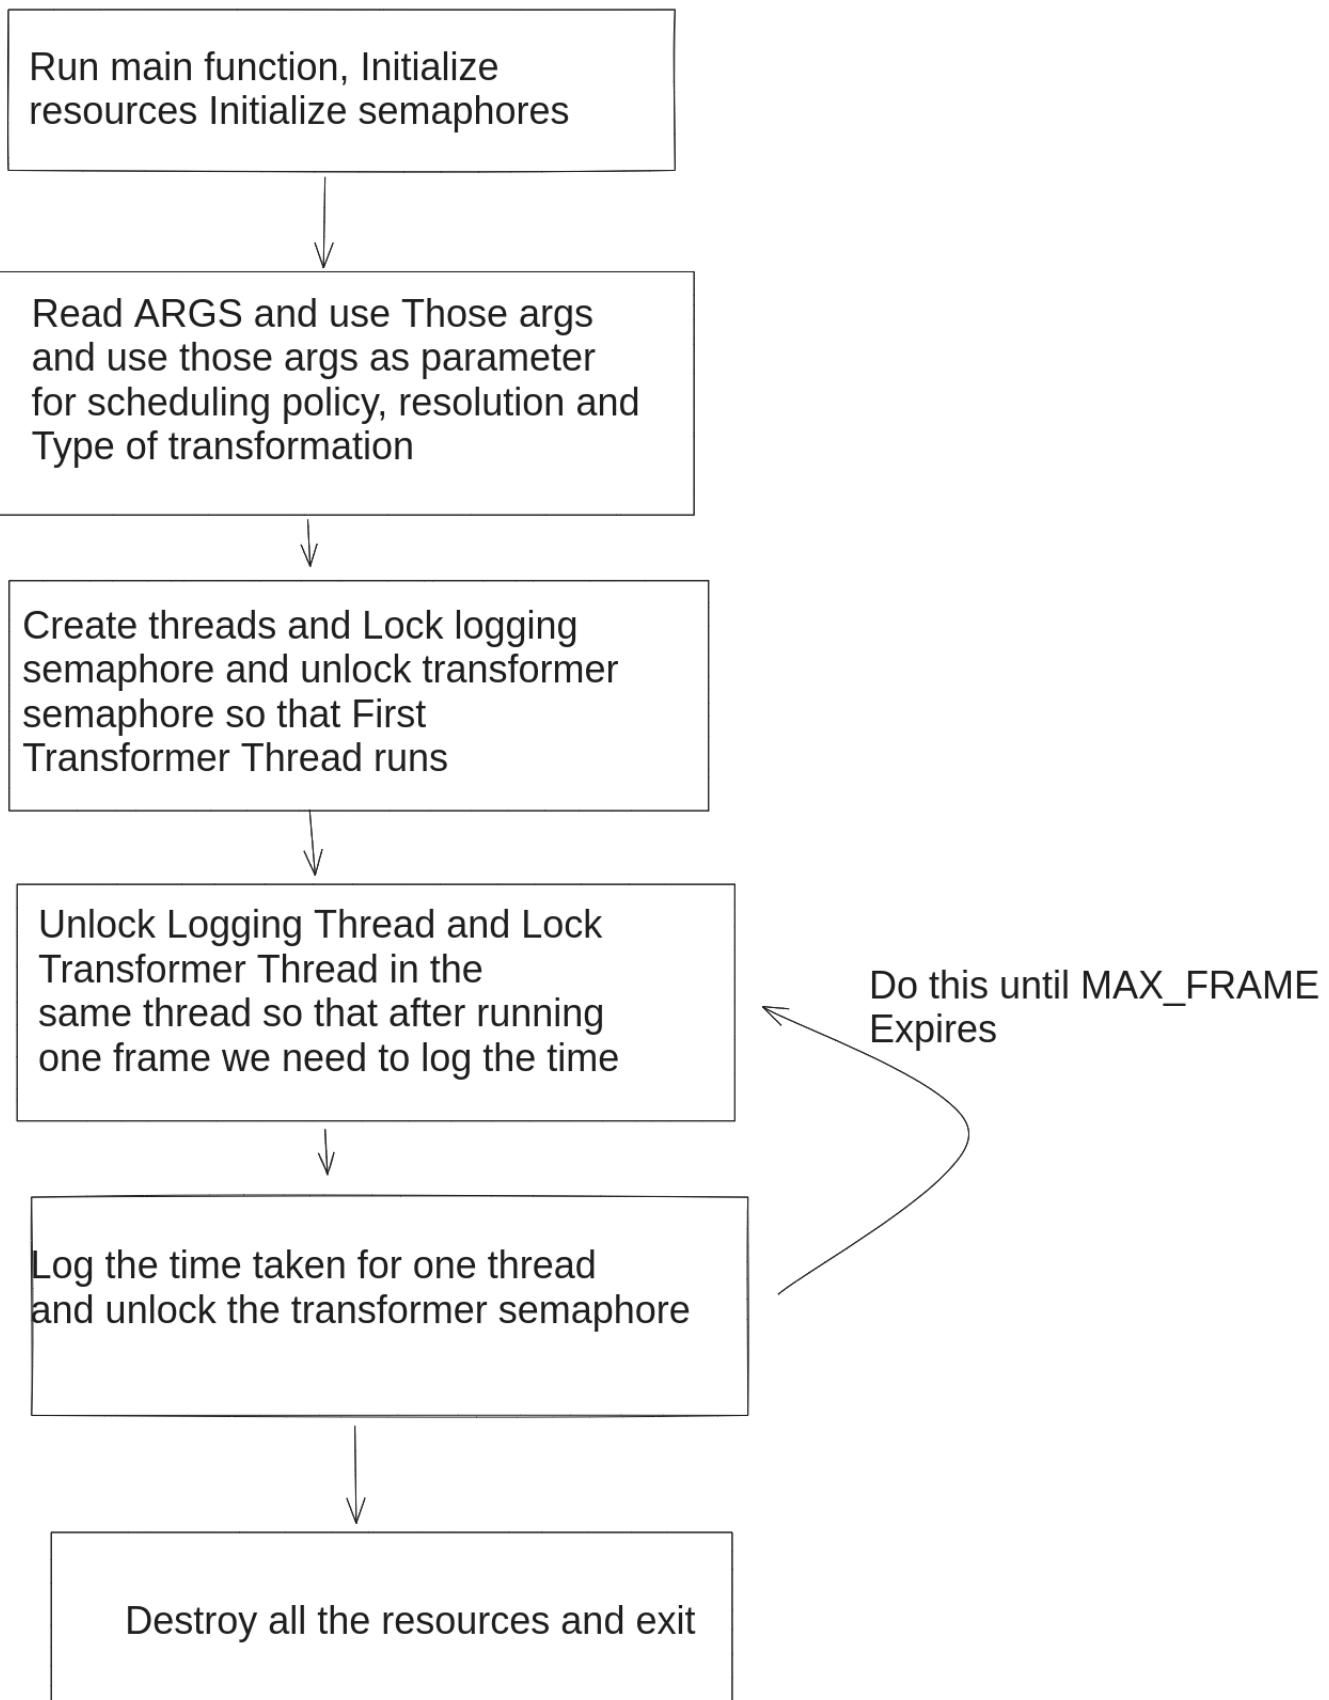


Figure 16: Flowchart

- Components to Include:
 - Video Capture and Processing: Utilizes OpenCV to capture video frames from a camera, apply image processing transformations (Canny edge detection, Hough Line and Circle

Transforms), and display the transformed images.

- Parallel Processing: Employs pthreads to create separate threads for image transformation and logging, allowing simultaneous processing and logging.
- Real-time Scheduling: Implements real-time scheduling policies (SCHED_FIFO, Sched opther) to prioritize the processing and logging tasks, aiming to meet soft real-time constraints.
- Synchronization: Uses semaphores to coordinate the execution between the transformation and logging threads, ensuring orderly processing and data logging.
- Logging and Performance Analysis: Records performance metrics (execution time, frame rate, scheduling policy, resolution, transformation type) to a CSV file for further analysis.

(b) Algorithm Analysis

- Transformation Descriptions: Describe each of the three chosen transformations in detail. For example, you might select grayscale conversion, edge detection, and Gaussian blur as your transformations. Explain the purpose of each transformation and its expected impact on processing time.

- Scheduling Approach:

- Semaphore Initialization

Two semaphores, transformation_semaphore and logging_semaphore, are initialized at the beginning of the program. These semaphores are used to synchronize the transformation of video frames (using different image processing techniques like Canny, Hough Line, or Hough Circle Transform) and the logging of performance metrics.

- Synchronization Between Processing and Logging

Transformation Semaphore (transformation_semaphore): This semaphore is used to control the execution flow between capturing/processing video frames and the logging thread. Initially, the semaphore is posted (incremented) to allow the transformation thread to start processing a frame.

Logging Semaphore (logging_semaphore): Conversely, this semaphore is waited on (decremented) by the logging thread to ensure it logs performance data only after a frame has been processed.

Execution Flow Video Frame Processing: The transformation thread starts by capturing a frame from the video feed. After processing the frame (applying the selected image transformation technique), the thread posts (increments) the logging_semaphore, signaling that the logging thread can now proceed to log the performance data for that frame.

Performance Logging: The logging thread, upon being allowed to proceed (after the logging_semaphore is posted by the transformation thread), calculates the processing time and logs the performance metrics. Once logging is complete, it posts (increments) the transformation_semaphore, indicating that the transformation thread can process the next frame.

Synchronization Loop: This process of alternately waiting on and posting semaphores creates a synchronized loop between the transformation and logging threads, ensuring that performance data for each frame is logged sequentially and accurately reflects the processing times.

Soft Exit Mechanism: The soft_exit flag is used as a soft mechanism to exit the loops in both transformation and logging threads safely. It ensures that both threads can complete their current iteration before closing, maintaining data integrity and avoiding premature termination that could lead to missed logging or incomplete frame processing.

(c) Output Screenshots

- Implementation Details: Briefly describe how the prototype was implemented, including the programming language, libraries (e.g., OpenCV for transformations), and tools used.
- Data Collection: Shows the resolution, scheduler policy, and the number of processors. The scheduler policy changes from SCHED_OTHER to SCHED_FIFO, indicating a switch to a real-time scheduling policy for more predictable execution.
- Execution Logs:
Each "working log" entry shows the frame processing time and calculates the frame rate as the inverse of this time.
Variability in execution time (jitter) is observable through the differences in "total_time taken for 1 frame" across frames.
- The Average FPS is calculated over the run, giving a measure of overall performance. The average execution time gives an idea of how long, on average, each frame takes to process.
- Total deadline miss: Indicates how many times the frame processing failed to meet the soft deadline, reflecting on the real-time performance and system load. For -transform=Canny

```

1 CSV header is correct.
2 Resolution: 640x480
3 Scheduler Policy: 1
4 This system has 12 processors configured and 12 processors available.
5 Before adjustments to scheduling policy:
6 Pthread Policy is SCHED_OTHER
7 After adjustments to scheduling policy:
8 Pthread Policy is SCHED_FIFO
9 Setting thread 0 to core 0
10 Setting thread 1 to core 0
11 Threads created
12 Joining thread 0
13 working log [ WARN:0] global ./modules/videoio/src/cap_gstreamer.cpp (1100)
14 open OpenCV | GStreamer warning: Cannot query video position: status=0,
15 value=-1, duration=-1
16 working logcurrent fps : 6.800981 , total_time taken for 1 frame: 0.147038
17 working logcurrent fps : 7.669834 , total_time taken for 1 frame: 0.130381
18 working logcurrent fps : 7.822449 , total_time taken for 1 frame: 0.127837
19 working logcurrent fps : 7.325486 , total_time taken for 1 frame: 0.136510
20 working logcurrent fps : 7.348487 , total_time taken for 1 frame: 0.136082
21 working logcurrent fps : 7.439587 , total_time taken for 1 frame: 0.134416
22 working logcurrent fps : 7.400848 , total_time taken for 1 frame: 0.135120
23 working logcurrent fps : 7.616983 , total_time taken for 1 frame: 0.131286
24 working logcurrent fps : 7.260292 , total_time taken for 1 frame: 0.137736
25 Thread 0 joined successfully.
26 Joining thread 1
27 Thread 1 joined successfully.
28 Average FPS 8.220948, Average time 0.121640
29 Total deadline miss : 5

```

For -transform=HoughLine

```

1 CSV header is correct.
2 Resolution: 640x480
3 Scheduler Policy: 1
4 This system has 12 processors configured and 12 processors available.
5 Before adjustments to scheduling policy:
6 Pthread Policy is SCHED_OTHER
7 After adjustments to scheduling policy:
8 Pthread Policy is SCHED_FIFO
9 Setting thread 0 to core 0
10 Setting thread 1 to core 0
11 WoringThreads created
12 Joining thread 0
13 working log [ WARN:0] global ./modules/videoio/src/cap_gstreamer.cpp (1100)
14 open OpenCV | GStreamer warning: Cannot query video position: status=0,
15 value=-1, duration=-1
16 working logcurrent fps : 5.931778 , total_time taken for 1 frame: 0.168584

```

```

17 working logcurrent fps : 9.412008 , total_time taken for 1 frame: 0.106247
18 working logcurrent fps : 7.601903 , total_time taken for 1 frame: 0.131546
19 working logcurrent fps : 7.402297 , total_time taken for 1 frame: 0.135093
20 working logcurrent fps : 7.084185 , total_time taken for 1 frame: 0.141160
21 working logcurrent fps : 7.610951 , total_time taken for 1 frame: 0.131390
22 working logcurrent fps : 7.393532 , total_time taken for 1 frame: 0.135253
23 working logcurrent fps : 7.787055 , total_time taken for 1 frame: 0.128418
24 working logcurrent fps : 7.366140 , total_time taken for 1 frame: 0.135756
25 Thread 0 joined successfully.
26 Joining thread 1
27 Thread 1 joined successfully.
28 Average FPS 8.240986, Average time 0.121345
29 Total deadline miss : 5
30
31

```

Fpr -transform=HoughCircle

```

1 CSV header is correct.
2 Resolution: 640x480
3 Scheduler Policy: 1
4 This system has 12 processors configured and 12 processors available.
5 Before adjustments to scheduling policy:
6 Pthread Policy is SCHED_OTHER
7 After adjustments to scheduling policy:
8 Pthread Policy is SCHED_FIFO
9 Setting thread 0 to core 0
10 Setting thread 1 to core 0
11 Threads created
12 Joining thread 0
13 working log [ WARN:0] global ./modules/videoio/src/cap_gstreamer.cpp (1100)
14 open OpenCV | GStreamer warning: Cannot query video position: status=0,
15 value=-1, duration=-1
16 working logcurrent fps : 7.336637 , total_time taken for 1 frame: 0.136302
17 working logcurrent fps : 6.928957 , total_time taken for 1 frame: 0.144322
18 working logcurrent fps : 8.009877 , total_time taken for 1 frame: 0.124846
19 working logcurrent fps : 7.141519 , total_time taken for 1 frame: 0.140026
20 working logcurrent fps : 7.693228 , total_time taken for 1 frame: 0.129984
21 working logcurrent fps : 7.618371 , total_time taken for 1 frame: 0.131262
22 working logcurrent fps : 7.320480 , total_time taken for 1 frame: 0.136603
23 working logcurrent fps : 7.581601 , total_time taken for 1 frame: 0.131898
24 working logcurrent fps : 7.249677 , total_time taken for 1 frame: 0.137937
25 Thread 0 joined successfully.
26 Joining thread 1
27 Thread 1 joined successfully.
28 Average FPS 8.242794, Average time 0.121318
29 Total deadline miss : 5
30

```

- Bash Script: This bash script is designed to automate the testing of a program (./program) with different sets of parameters: transformation types, resolutions, and scheduler policies. It systematically iterates through each combination of these parameters, executes the program with them, and checks for any errors during execution.
- Python Script: This Python script is designed to load data from a CSV file, process and analyze this data, and finally plot specific portions of the data, focusing on execution times. It uses the pandas library for data manipulation and matplotlib for plotting

(d) Final Predictable Response Jitter Analysis

- Jitter Analysis Methodology: Describe how jitter was calculated, considering the variation in frame processing times relative to the set deadline. Explain the statistical methods used to analyze jitter.
- Graphical Representation:
 - Graph 1 (Jitter vs. Number of Frames):

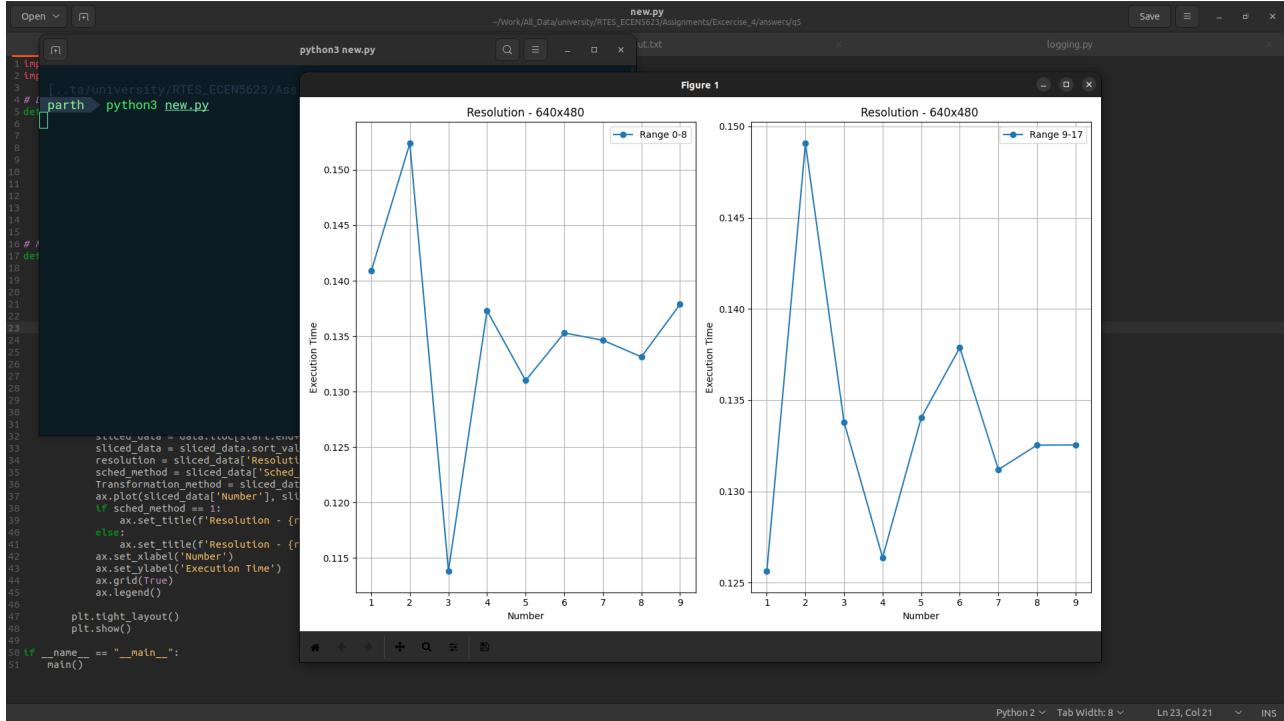


Figure 17: Graph 1

– Graph 2 (Number of Frames & Execution Time):

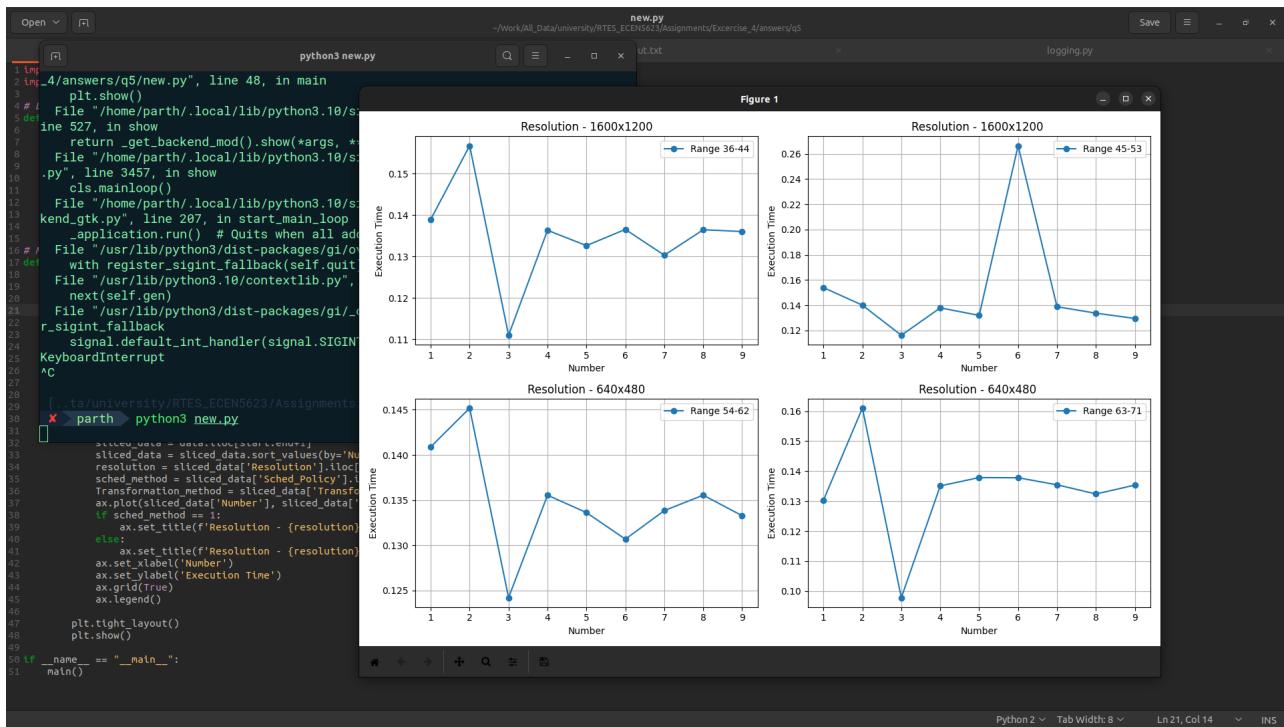


Figure 18: Graph 2

– Graph 2 (Number of Frames & Execution Time):

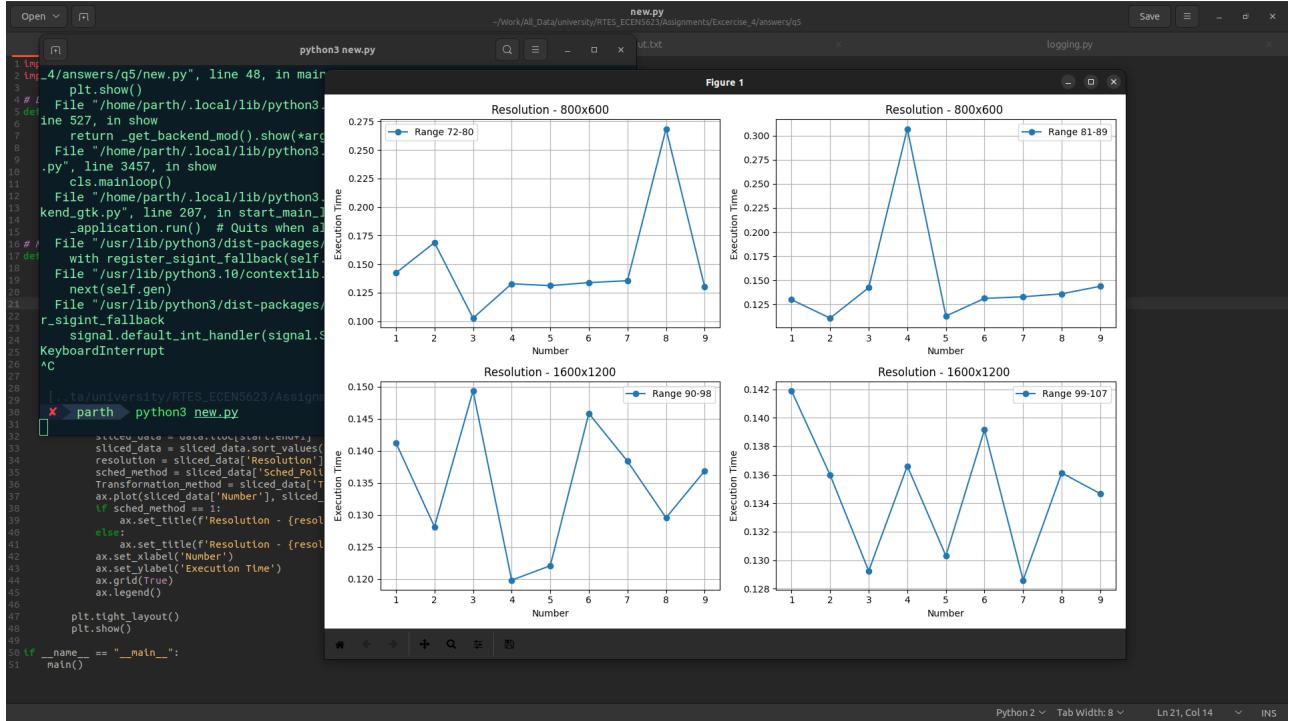


Figure 19: Graph 3

– Graph 2 (Number of Frames & Execution Time):

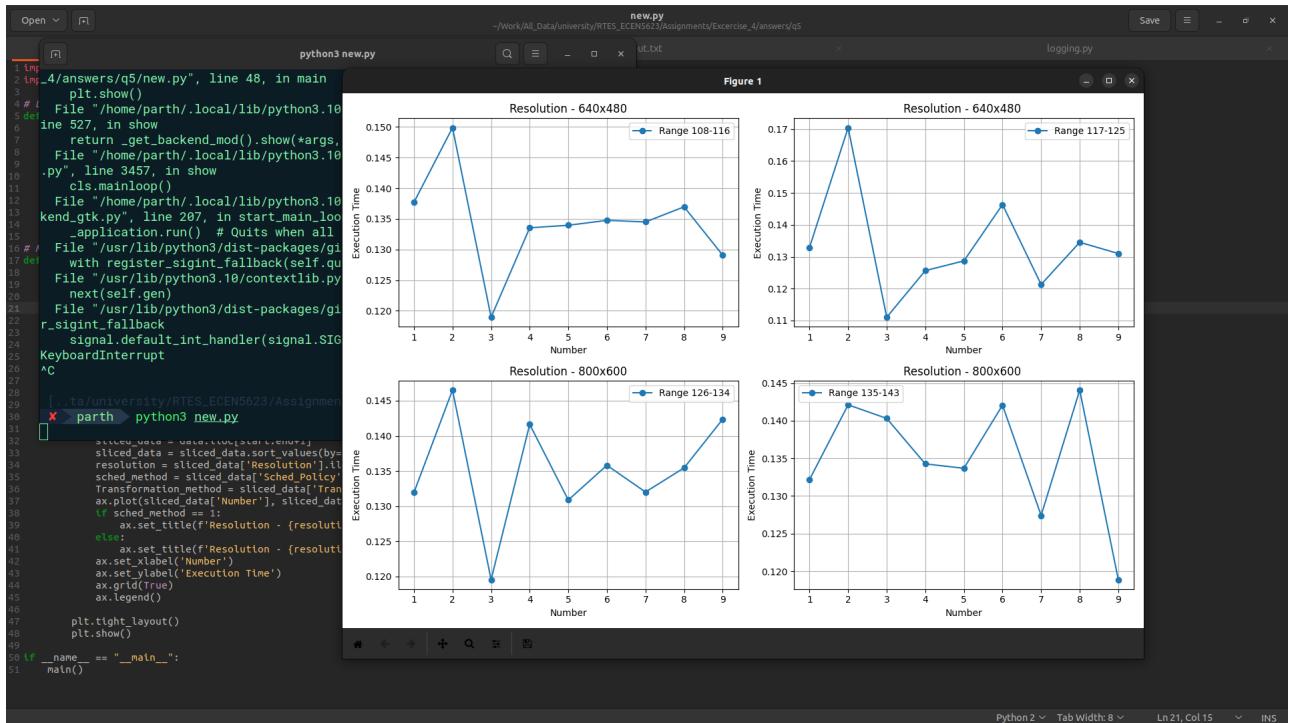


Figure 20: Graph 4

6 Question 6

Q: [10 points] Demonstrate the results and answer questions from the TA regarding the code you developed in #2, #3, #4, and #5.

7 References

1. ECEN 5623 Lecture slides material and example codes.
2. REAL-TIME EMBEDDED COMPONENTS AND SYSTEMS with LINUX and RTOS, Sam Siewert John Pratt (Chapter 6, 7 & 8).
3. Exercise 4 requirements included links and documentation.

Appendices

A C/Cpp/Python Code for the Implementation

Ex_3_script.sh

```
1 #!/bin/bash
2
3 matrix=("Canny" "HoughLine" "HoughCircle")
4 resolutions=("640x480" "800x600" "1600x1200")
5 sched_policies=("SCHED_FIFO" "SCHED_OTHER")
6
7 rm example.csv
8 touch example.csv
9
10 # Iterate through each combination of parameters
11 for mat in "${matrix[@]}"; do
12     for resolution in "${resolutions[@]}"; do
13         for policy in "${sched_policies[@]}"; do
14             echo "Running: ./program --transformation=${mat} --resolution=
${resolution} --sched_policy=${policy}"
15             sudo ./program --transformation=${mat} --resolution=${resolution} --
16             sched_policy=${policy}
17             return_value="$?"
18             # Check if the program exited with an error
19             if [ "$return_value" -ne 0 ]; then
20                 echo "Error running program with parameters --transformation=${mat} --
21                 resolution=${resolution} --sched_policy=${policy}, exiting."
22                 exit -1
23             fi
24         done
25     done
done
```

```
example.py

1 import matplotlib.pyplot as plt
2 import pandas as pd
3
4 # Load data from CSV file
5 def load_data_from_csv(filename):
6     try:
7         df = pd.read_csv(filename)
8         return df
9     except FileNotFoundError:
10        print(f"File {filename} not found.")
11        return None
12    except Exception as e:
13        print(f"An error occurred while loading data: {e}")
14        return None
15
16 # Plot data
17 def plot_data(dataframe):
18     if dataframe is not None:
19         x = dataframe.iloc[:, 0] # Assuming first column is x-axis data
20         y = dataframe.iloc[:, 1] # Assuming second column is y-axis data
21
22         plt.plot(x, y)
23         plt.xlabel('X Label')
24         plt.ylabel('Y Label')
25         plt.title('Data Plot')
26         plt.grid(True)
27         plt.show()
28
29 # Main function
30 def main():
31     csv_filename = 'example.csv'
32     data = load_data_from_csv(csv_filename)
33
34     # Grouping the data by 'Sched_Policy' and 'Transform', then calculate standard
35     # deviation
36     grouped_std_dev = data.groupby(['Sched_Policy', 'Transform', 'Resolution'])['
37     Execution_time'].std().reset_index()
38     print(grouped_std_dev);
39
40     grouped_std_dev
41
42     # if data is not None:
43     #     plot_data(data)
44
45 if __name__ == "__main__":
46     main()
```

program.cpp

```
1 #include "opencv2/imgcodecs.hpp"
2 #include "opencv2/highgui.hpp"
3 #include "opencv2/imgproc.hpp"
4 #include <iostream>
5 #include <string>
6 #include <map>
7 #include <cstdlib> // For std::exit and EXIT_FAILURE
8 #include <pthread.h>
9 #include <sched.h>
10 #include <time.h>
11 #include <fstream>
12
13 #include <pthread.h>
14 #include <stdio.h>
15 #include <time.h>
16 #include <aio.h>
17 #include <math.h>
18 #include <unistd.h>
19 #include <errno.h>
20 #include <stdlib.h>
21 #include <sched.h>
22 #include <stddef.h>
23 #include <sys/sysinfo.h>
24 #include <semaphore.h>
25
26 #define MAX_FRAME 10
27 #define NUM_OF_THREAD 2
28 #define ESCAPE_KEY (27)
29 #define SYSTEM_ERROR (-1)
30 #define TARGET_CORE 0
31
32 #define SOFT_DEADLINE 0.135 // ms
33
34 sem_t transformation_semaphore, logging_semaphore;
35 static int total_deadline_miss = 0;
36 bool initial = true;
37
38 void *(*transformationType)(void *);
39 std::string resolution;
40 int schedPolicy = SCHED_FIFO; // Default policy
41 std::string transformationStr;
42 volatile int soft_exit = 0;
43 double sum_of_execution = 0;
44
45 std::string filePath = "./example.csv";
46
47 struct timespec update_interval;
48 struct timespec last_update_interval;
49 struct timespec start_interval;
50 struct timespec end_interval;
51
52 namespace canny
53 {
```

```

54     int lowThreshold = 0;
55     const int max_lowThreshold = 100;
56     const int ratio = 3;
57     const int kernel_size = 3;
58     const char *window_name = "Edge Map";
59 }
60
61 typedef struct
62 {
63     int threadId;
64     int width;
65     int height;
66     cv::Mat frame;
67     cv::Mat intermediate_src;
68     cv::Mat output;
69     std::string window_name;
70
71 } ThreadArgs_t;
72
73 typedef struct
74 {
75     int period;
76     int burst_time;
77     struct sched_param priority_param;
78     void *(*thread_handle)(void *);
79     pthread_t thread;
80     ThreadArgs_t thread_args;
81     void *return_Value;
82     pthread_attr_t attribute;
83     int target_cpu;
84 } RmTask_t;
85
86 // Mat frame;
87 void print_scheduler(void)
88 {
89     int schedType;
90     schedType = sched_getscheduler(getpid());
91     switch (schedType)
92     {
93         case SCHED_FIFO:
94             printf("Pthread Policy is SCHED_FIFO\n");
95             break;
96         case SCHED_OTHER:
97             printf("Pthread Policy is SCHED_OTHER\n");
98             break;
99         case SCHED_RR:
100            printf("Pthread Policy is SCHED_OTHER\n");
101            break;
102        default:
103            printf("Pthread Policy is UNKNOWN\n");
104    }
105 }
106
107 void appendDataToCsv(int max_frame, float total_time)
108 {
109     // Open the file in append mode

```

3/24/24, 5:19 PM

program.cpp

```
110     std::ofstream file(filePath, std::ios::app);
111
112     if (!file.is_open())
113     {
114         std::cerr << "Failed to open the file for appending." << std::endl;
115         return;
116     }
117
118     // Write the new data to the file
119     file << MAX_FRAME - max_frame << "," << total_time << "," << (1 / total_time) <<
120     ", " << schedPolicy << "," << resolution << "," << transformationStr << std::endl;
121
122     // Close the file
123     file.close();
124 }
125
126 void HoughLinePTransform(RmTask_t *data)
127 {
128
129     std::vector<cv::Vec4i> linesP;
130     // Will hold the results of the detection
131     cv::HoughLinesP(data->thread_args.intermediate_src, linesP, 1, CV_PI / 180, 50,
132     50, 10); // Runs the actual detection
133
134     // Draw the lines
135     for (size_t i = 0; i < linesP.size(); i++)
136     {
137         cv::Vec4i l = linesP[i];
138         cv::line(data->thread_args.output, cv::Point(l[0], l[1]), cv::Point(l[2],
139         l[3]), cv::Scalar(0, 0, 255), 3, cv::LINE_AA);
140         cv::line(data->thread_args.frame, cv::Point(l[0], l[1]), cv::Point(l[2], l[3]),
141         cv::Scalar(0, 0, 255), 3, cv::LINE_AA);
142     }
143     cv::imshow(data->thread_args.window_name, data->thread_args.output);
144 }
145
146 // Function to detect and draw circles in an image using the Hough Transform
147 void HoughCircleCam(RmTask_t *data)
148 {
149
150     cvtColor(data->thread_args.frame, data->thread_args.intermediate_src,
151     cv::COLOR_BGR2GRAY);
152     cv::medianBlur(data->thread_args.intermediate_src, data->
153     thread_args.intermediate_src, 5);
154     std::vector<cv::Vec3f> circles;
155
156     cv::HoughCircles(data->thread_args.intermediate_src, circles, cv::HOUGH_GRADIENT,
157     1,
158             data->thread_args.intermediate_src.rows / 16, // Change this
159             value to detect circles with different distances to each other
160             100, 30, 1, 30 // Change the last
161             two parameters (min_radius & max_radius) to detect larger circles
162     );
163
164     for (size_t i = 0; i < circles.size(); i++)
165     {
166         cv::Vec3i c = circles[i];
167         cv::Point center = cv::Point(c[0], c[1]); // Circle center
```

3/24/24, 5:19 PM

program.cpp

```
158     circle(data->thread_args.frame, center, 1, cv::Scalar(0, 100, 100), 3,
159     cv::LINE_AA);
160     int radius = c[2];
161     circle(data->thread_args.frame, center, radius, cv::Scalar(255, 0, 255), 3,
162     cv::LINE_AA);
163 }
164
165
166 void CannyThreshold(RmTask_t *data)
167 {
168     cvtColor(data->thread_args.frame, data->thread_args.intermediate_src,
169     cv::COLOR_BGR2GRAY);
170
171     /// Reduce noise with a kernel 3x3
172     cv::blur(data->thread_args.intermediate_src, data->thread_args.output,
173     cv::Size(3, 3));
174
175     /// Canny detector
176     cv::Canny(data->thread_args.output, data->thread_args.output,
177     canny::lowThreshold, canny::lowThreshold * canny::ratio, canny::kernel_size);
178
179     // Using Canny's output as a mask, we display our result
180     cv::Mat timg_grad;
181
182     // Initialize `timg_grad` to be the same size and type as the source frame, but
183     // filled with zeros
184     // Assuming `data->thread_args.frame` is a cv::Mat
185     timg_grad = cv::Mat::zeros(data->thread_args.frame.size(), data->
186     thread_args.frame.type());
187
188     // Using Canny's output as a mask, copy the frame to `timg_grad` wherever the
189     // mask is non-zero
190     data->thread_args.frame.copyTo(timg_grad, data->thread_args.output);
191
192     // Display the result
193     cv::imshow(data->thread_args.window_name, timg_grad);
194 }
195
196 void *HoughLineTransform(void *args)
197 {
198     printf("Woring");
199     RmTask_t *data = static_cast<RmTask_t *>(args);
200     data->thread_args.window_name = "Houghline Transformation";
201     cv::VideoCapture cam0(0);
202     cv::namedWindow("video_display");
203     char winInput;
204
205     if (!cam0.isOpened())
206     {
207         exit(SYSTEM_ERROR);
208     }
209
210     cam0.set(cv::CAP_PROP_FRAME_WIDTH, data->thread_args.width);
211     cam0.set(cv::CAP_PROP_FRAME_HEIGHT, data->thread_args.height);
212
213     int max_frame = MAX_FRAME;
```

3/24/24, 5:19 PM

program.cpp

```
208     clock_gettime(CLOCK_REALTIME, &start_interval);
209
210     while (max_frame)
211     {
212         sem_wait(&transformation_semaphore);
213         cam0.read(data->thread_args.frame);
214         cv::imshow("video_display", data->thread_args.frame);
215         cv::Canny(data->thread_args.frame, data->thread_args.intermediate_src, 80,
216 240, 3);
217         cv::cvtColor(data->thread_args.intermediate_src, data->thread_args.output,
218 cv::COLOR_GRAY2BGR);
219         cv::namedWindow(data->thread_args.window_name, cv::WINDOW_AUTOSIZE);
220         HoughLinePTransform(data);
221         if ((winInput = cv::waitKey(10)) == ESCAPE_KEY)
222         {
223             soft_exit = 1;
224             sem_post(&logging_semaphore);
225             cv::destroyAllWindows();
226             break;
227         }
228     }
229
230     return nullptr;
231 }
232
233 void *HoughCircleTransform(void *args)
234 {
235
236     RmTask_t *data = static_cast<RmTask_t *>(args);
237     data->thread_args.window_name = "HoughCircle Transformation";
238     cv::VideoCapture cam0(0);
239     cv::namedWindow("video_display");
240     char winInput;
241     if (!cam0.isOpened())
242     {
243         exit(SYSTEM_ERROR);
244     }
245     cam0.set(cv::CAP_PROP_FRAME_WIDTH, data->thread_args.width);
246     cam0.set(cv::CAP_PROP_FRAME_HEIGHT, data->thread_args.height);
247
248     int max_frame = MAX_FRAME;
249     clock_gettime(CLOCK_REALTIME, &start_interval);
250     while (max_frame)
251     {
252         sem_wait(&transformation_semaphore);
253         cam0.read(data->thread_args.frame);
254         cv::imshow("video_display", data->thread_args.frame);
255         HoughCircleCam(data);
256         if ((winInput = cv::waitKey(10)) == ESCAPE_KEY)
257         {
258             soft_exit = 1;
259             sem_post(&logging_semaphore);
260             cv::destroyAllWindows();
261             break;
262         }
263     }
264 }
```

```
263     else if (winInput == 'n')
264     {
265         printf("input %c is ignored\n", winInput);
266     }
267     max_frame--;
268     sem_post(&logging_semaphore);
269 }
270
271 return nullptr;
272 }
273
274 void *CannyTransform(void *args)
275 {
276
277     RmTask_t *data = static_cast<RmTask_t *>(args);
278     data->thread_args.window_name = "Canny Transformation";
279     cv::VideoCapture cam0(0);
280     cv::namedWindow("video_display");
281     char winInput;
282
283     if (!cam0.isOpened())
284     {
285         exit(SYSTEM_ERROR);
286     }
287
288     cam0.set(cv::CAP_PROP_FRAME_WIDTH, data->thread_args.width);
289     cam0.set(cv::CAP_PROP_FRAME_HEIGHT, data->thread_args.height);
290
291     int max_frame = MAX_FRAME;
292
293     clock_gettime(CLOCK_REALTIME, &start_interval);
294     while (max_frame)
295     {
296         sem_wait(&transformation_semaphore);
297
298         cam0.read(data->thread_args.frame);
299
300         cv::imshow("Canny Video", data->thread_args.frame);
301
302         cv::namedWindow(data->thread_args.window_name, cv::WINDOW_AUTOSIZE);
303
304         CannyThreshold(data);
305
306         if ((winInput = cv::waitKey(10)) == ESCAPE_KEY)
307         {
308             soft_exit = 1;
309             sem_post(&logging_semaphore);
310             cv::destroyAllWindows();
311             break;
312         }
313
314         max_frame--;
315         sem_post(&logging_semaphore);
316     }
317
318     return nullptr;
319 }
```

```

319 }
320
321 void *LoggingThread(void *args)
322 {
323     int max_frame = MAX_FRAME;
324     clock_gettime(CLOCK_REALTIME, &last_update_interval);
325     while (max_frame)
326     {
327         printf("working log");
328
329         sem_wait(&logging_semaphore);
330
331         clock_gettime(CLOCK_REALTIME, &update_interval);
332         float total_sec = update_interval.tv_sec - last_update_interval.tv_sec;
333         float total_ns = ((float)(update_interval.tv_nsec -
last_update_interval.tv_nsec) / 1000000000);
334         float total_time = total_sec + total_ns;
335
336         if (!initial)
337         {
338
339             if (total_time > SOFT_DEADLINE)
340             {
341                 total_deadline_miss++;
342             }
343             sum_of_execution += total_time;
344
345             printf("current fps/jitter: update_interval - last_update_interval : %f ,
total_time taken for 1 frame: %f \n\r", (1 / total_time), total_time);
346
347             appendDataToCsv(max_frame, total_time);
348
349         }
350
351         if (soft_exit)
352         {
353             printf("Exiting thread 2\n\r");
354             break;
355         }
356         sem_post(&transformation_semaphore);
357         last_update_interval = update_interval;
358         max_frame--;
359         initial = false;
360     }
361     return nullptr;
362 }
363
364 int main(int argc, char **argv)
365 {
366
367     sem_init(&transformation_semaphore, false, 0);
368     sem_init(&logging_semaphore, false, 0);
369
370     sem_post(&transformation_semaphore);
371
372     const std::string expectedHeader = "Number,Execution_time,Frame_rate,
Sched_Policy,Resolution,Transform";

```

```
373     std::ifstream file(filePath);
374     std::string currentHeader;
375
376     if (file.is_open())
377     {
378         // Get the first line from the file
379         std::getline(file, currentHeader);
380         file.close();
381
382         // Check if the current header matches the expected header
383         if (currentHeader != expectedHeader)
384         {
385             // The headers do not match, so we need to write the correct header
386
387             // Store the rest of the file content
388             std::string restOfFileContent;
389             std::string line;
390             while (std::getline(file, line))
391             {
392                 restOfFileContent += line + "\n";
393             }
394
395             // Write the correct header and the rest of the file
396             std::ofstream outFile(filePath);
397             outFile << expectedHeader << "\n"
398                 << restOfFileContent;
399             outFile.close();
400
401             std::cout << "Header was corrected in the CSV file." << std::endl;
402         }
403     }
404     else
405     {
406         std::cout << "CSV header is correct." << std::endl;
407     }
408 }
409 else
410 {
411     std::cerr << "Could not open the file." << std::endl;
412     return 1;
413 }
414
415 void *(*transformationType) (void *) = CannyTransform;
416 std::map<std::string, int> schedPolicyMap = {
417     {"SCHED_FIFO", SCHED_FIFO},
418     {"SCHED_RR", SCHED_RR},
419     {"SCHED_OTHER", SCHED_OTHER}};
420
421 std::map<std::string, void *(&)(void *)> TransformationPolicyMap = {
422     {"Canny", CannyTransform},
423     {"HoughLine", HoughLineTransform},
424     {"HoughCircle", HoughCircleTransform}};
425
426 for (int i = 1; i < argc; i++)
427 {
428     std::string arg = argv[i];
```

3/24/24, 5:19 PM

program.cpp

```
429     if (arg.find("--transformation=") == 0)
430     {
431         transformationStr = arg.substr(std::string("--transformation=").length())
432     ;
433         if (TransformationPolicyMap.find(transformationStr) != TransformationPolicyMap.end())
434         {
435             transformationType = TransformationPolicyMap[transformationStr];
436         }
437         else
438         {
439             std::cerr << "Unknown transformation type: " << transformationStr <<
440             std::endl;
441             std::exit(EXIT_FAILURE);
442         }
443     }
444     else if (arg.find("--resolution=") == 0)
445     {
446         resolution = arg.substr(std::string("--resolution=").length());
447     }
448     else if (arg.find("--sched_policy=") == 0)
449     {
450         std::string policyStr = arg.substr(std::string("--sched_policy=").length());
451         if (schedPolicyMap.find(policyStr) != schedPolicyMap.end())
452         {
453             schedPolicy = schedPolicyMap[policyStr];
454         }
455         else
456         {
457             std::cerr << "Unknown scheduler policy: " << policyStr << std::endl;
458             std::exit(EXIT_FAILURE);
459         }
460     }
461     std::cout << "Resolution: " << resolution << std::endl;
462     std::cout << "Scheduler Policy: " << schedPolicy << std::endl;
463
464     pthread_t threads[NUM_OF_THREAD];
465     cpu_set_t threadcpu;
466
467     CPU_SET(TARGET_CORE, &threadcpu);
468
469     RmTask_t tasks[NUM_OF_THREAD] = {
470         .period = 20,      // ms
471         .burst_time = 10, // ms
472         .priority_param = {0},
473         .thread_handle = transformationType,
474         .thread = threads[0],
475         .thread_args = {0, 0, 0},
476         .return_Value = NULL,
477         .attribute = {0, 0},
478         .target_cpu = TARGET_CORE},
479
480         {.period = 20,      // ms
481         .burst_time = 10, // ms
```

```
482     .priority_param = {0},  
483     .thread_handle = LoggingThread,  
484     .thread = threads[0],  
485     .thread_args = {0, 0, 0},  
486     .return_Value = NULL,  
487     .attribute = {0, 0},  
488     .target_cpu = TARGET_CORE},  
489 };  
490  
491 pthread_attr_t attribute_flags_for_main; // for schedular type, priority  
492 struct sched_param main_priority_param;  
493  
494 printf("This system has %d processors configured and %d processors available.\n",  
get_nprocs_conf(), get_nprocs());  
495  
496 printf("Before adjustments to scheduling policy:\n");  
497 print_scheduler();  
498  
499 int rt_max_prio = sched_get_priority_max(schedPolicy);  
500 int rt_min_prio = sched_get_priority_min(schedPolicy);  
501  
502 main_priority_param.sched_priority = rt_max_prio;  
503 for (int i = 0; i < NUM_OF_THREAD; i++)  
504 {  
505     tasks[i].priority_param.sched_priority = rt_max_prio;  
506  
507     // initialize attributes  
508     pthread_attr_init(&tasks[i].attribute);  
509  
510     pthread_attr_setinheritsched(&tasks[i].attribute, PTHREAD_EXPLICIT_SCHED);  
511     pthread_attr_setschedpolicy(&tasks[i].attribute, schedPolicy);  
512     pthread_attr_setschedparam(&tasks[i].attribute, &tasks[i].priority_param);  
513     pthread_attr_setaffinity_np(&tasks[i].attribute, sizeof(cpu_set_t), &  
threadcpu);  
514 }  
515  
516     pthread_attr_init(&attribute_flags_for_main);  
517  
518 // pthread_attr_setinheritsched(&attribute_flags_for_main,  
PTHREAD_EXPLICIT_SCHED);  
519     pthread_attr_setschedpolicy(&attribute_flags_for_main, schedPolicy);  
520     pthread_attr_setaffinity_np(&attribute_flags_for_main, sizeof(cpu_set_t), &  
threadcpu);  
521  
522     // Main thread is already created we have to modify the priority and scheduling  
scheme  
523     int status_setting_scheduler = sched_setscheduler(getpid(), schedPolicy, &  
main_priority_param);  
524     if (status_setting_scheduler)  
525     {  
526         printf("ERROR: sched_setscheduler rc is %d\n", status_setting_scheduler);  
527         perror(NULL);  
528         exit(-1);  
529     }  
530  
531     printf("After adjustments to scheduling policy:\n");  
532     print_scheduler();
```

```

533     int width = 0, height = 0;
534     size_t xPosition = resolution.find('x');
535     if (xPosition != std::string::npos)
536     {
537         width = std::stoi(resolution.substr(0, xPosition));
538         height = std::stoi(resolution.substr(xPosition + 1));
539     }
540 }
541
542 for (int i = 0; i < NUM_OF_THREAD; i++)
543 {
544     // Create a thread
545     // First parameter is thread which we want to create
546     // Second parameter is the flags that we want to give it to
547     // third parameter is the routine we want to give
548     // Fourth parameter is the value
549     printf("Setting thread %d to core %d\n", i, TARGET_CORE);
550
551     if (pthread_create(&tasks[i].thread, &tasks[i].attribute, tasks[i]
552 .thread_handle, &tasks[i]) != 0)
553     {
554         perror("Create_Fail");
555     }
556     printf("Threads created \n\r");
557     for (int i = 0; i < NUM_OF_THREAD; i++)
558     {
559         printf("Joining thread %d\n", i);
560         int join_result = pthread_join(tasks[i].thread, &tasks[i].return_Value);
561         if (join_result != 0)
562         {
563             printf("pthread_join failed for thread %d: %s\n", i,
564 strerror(join_result));
565         }
566         else
567         {
568             printf("Thread %d joined successfully.\n", i);
569         }
570     }
571     if (pthread_attr_destroy(&tasks[0].attribute) != 0)
572         perror("attr destroy");
573     if (pthread_attr_destroy(&tasks[1].attribute) != 0)
574         perror("attr destroy");
575
576     clock_gettime(CLOCK_REALTIME, &end_interval);
577
578     float total_sec = end_interval.tv_sec - start_interval.tv_sec;
579     float total_ns = ((float)(end_interval.tv_nsec - start_interval.tv_nsec) /
580 1000000000.0);
581     float total_time = (total_sec + total_ns);
582
583     printf("Average FPS %f, Average time %f \n\r", (MAX_FRAME / sum_of_execution),
584 sum_of_execution/MAX_FRAME);
585     printf("Total deadline miss : %d\n\r", total_deadline_miss);
586     sem_destroy(&transformation_semaphore);
587     sem_destroy(&logging_semaphore);

```

3/24/24, 5:19 PM

program.cpp

```
586 |
587 |     return 0;
588 }
```

