

ECEN 5623, Real-Time Systems:

Exercise #3 – Threading/Tasking and Real-Time Synchronization

DUE: As Indicated on Canvas and in class

Please thoroughly read Chapters 6, 7 & 8 in the text.


Please see example code provided on Canvas, or at [Linux](#), [FreeRTOS](#), and [VxWorks](#).
Linux start code and example code can be found here - [Linux/code/](#)

Please complete this lab on an Altera DE1-SoC, Raspberry Pi or Jetson and provide evidence that you did your work with screenshots.

Exercise #3 Requirements:

- 1) [10 points] [All papers here also on Canvas] Read Sha, Rajkumar, et al paper, "[Priority Inheritance Protocols: An Approach to Real-Time Synchronization](#)"
 - a) Summarize 3 main key points the paper makes. Read [Dr. Siewert's summary paper on the topic as well](#) which might be easier to understand.
 - b) Read the historical positions of [Linus Torvalds as described by Jonathan Corbet](#) and [Ingo Molnar and Thomas Gleixner](#) on this topic as well. Take a position on this topic yourself and write at least one well supported paragraph or more to defend your position based on what we have learned in class. Does the PI-futex ([Futex](#), [Futexes are Tricky](#)) that is described by Ingo Molnar provide safe and accurate protection from un-bounded priority inversion as described in the paper? If not, what is different about it?
 - c) Note that more recently Red Hat has added support for priority ceiling emulation and priority inheritance and has a great overview on use of these POSIX real-time extension features here – [general real-time overview](#). The key systems calls are [pthread_mutexattr_setprotocol](#), [pthread_mutexattr_getprotocol](#) as well as [pthread_mutex_setpriorityceiling](#) and [pthread_mutex_getpriorityceiling](#). Why did some of the Linux community resist addition of priority ceiling and/or priority inheritance until more recently? (Linux has been distributed since the early 1990's and the problem and solutions were known before this).
 - d) Given that there are two solutions to unbounded priority inversion, priority ceiling emulation and priority inheritance, based upon your reading on support for both by POSIX real-time extensions, which do you think is best to use and why?

- 2) [25 points] Review the terminology guide (glossary in the textbook)
- Describe clearly what it means to write "thread safe" functions that are "re-entrant".
There are generally three main ways to do this: 1) pure functions that use only stack and have no global memory, 2) functions which use thread indexed global data, and 3) functions which use shared memory global data, but synchronize access to it using a MUTEX semaphore critical section wrapper.
 - Describe each method and how you would code it and how it would impact real-time threads/tasks.
 - Now, using a MUTEX, provide an example using RT-Linux Pthreads that does a thread safe update of a complex state (3 or more numbers – e.g., Latitude, Longitude and Altitude of a location) with a timestamp ([pthread_mutex_lock](#)). Your code should include two threads and one should update a timespec structure contained in a structure that includes a double precision attitude state of {Lat, Long, Altitude and Roll, Pitch, Yaw at Sample_Time} and the other should read it and never disagree on the values as function of time. You can just make up values for the navigational state using math library function generators (e.g., use simple periodic functions for Roll, Pitch, Yaw $\sin(x)$, $\cos(x^2)$, and $\cos(x)$, where x =time and linear functions for Lat, Long, Alt) and see http://linux.die.net/man/3/clock_gettime for how to get a precision timestamp). The second thread should read the times-stamped state without the possibility of data corruption (partial update of one of the 6 floating point values). There should be no disagreement between the functions and the state reader for any point in time. Run this for 180 seconds with a 1 Hz update rate and a 0.1 Hz read rate. Make sure the 18 values read are correct.



E.g., Function Generators - <https://www.desmos.com/calculator>

- 3) [20 points] Download [example-sync-updated-2/](#) and review, build, and run it.
- Describe both the issues of deadlock and unbounded priority inversion and the root cause for both in the example code.
 - Fix the deadlock so that it does not occur by using a random back-off scheme to resolve. For the unbounded inversion, is there a real fix in Linux – if not, why not?

- c) What about a patch for the Linux kernel? For example, Linux Kernel.org recommends the [RT PREEMPT Patch](#), also discussed by the [Linux Foundation Realtime Start](#) and [this blog](#), but would this really help? Read about the patch and describe why think it would or would not help with unbounded priority inversion. Based on inversion, does it make sense to simply switch to an RTOS and not use Linux at all for both HRT and SRT services?
- 4) [15 points] Review [POSIX-examples](#) and especially [POSIX MQ LOOP](#) and build the code related to POSIX message queues and run them to learn about basic use.
 - a) First, re-write the simple message queue demonstration code in `heap_mq.c` and `posix_mq.c` so that it uses RT-Linux Pthreads (FIFO) instead of `SCHEDULE_OTHER`, and then write a brief paragraph describing how the two message queue applications are similar and how they are different. Prove that you got the POSIX message queue features working in Linux on your target board.
 - b) Message queues are often suggested as a way to avoid MUTEX priority inversion. Would you agree that this really circumvents the problem of unbounded inversion? If so why, if not, why not?
- 5) [20 points] Watchdog timers, timeouts and timer services – First, read this overview of the [Linux Watchdog Daemon](#) and the Linux manual page on the watchdog daemon - <https://linux.die.net/man/8/watchdog> . Also see the [Watchdog Explained](#).
 - a) Describe how it might be used if software caused an indefinite deadlock.
 - b) Next, to explore timeouts, use your code from #2 and create a thread that waits on a MUTEX semaphore for up to 10 seconds and then un-blocks and prints out “No new data available at <time>” and then loops back to wait for a data update again. Use a variant of the [pthread_mutex_lock](#) called [pthread_mutex_timedlock](#) to solve this programming problem.
- 6) [10 points] Demonstrate the results and answer questions from the TA regarding the code you developed in #2, #3, #4, and #5.

Overall, provide a well-documented professional report of your findings, output, and tests so that it is easy for a colleague (or instructor) to understand what you’ve done. Include any C/C++ source code you write (or modify) and Makefiles needed to build your code. I will look at your report first, so it must be well written and clearly address each problem providing clear and concise responses to receive credit.

Note: Linux manual pages can be found for all system calls (e.g. `fork()`) on the web at <http://linux.die.net/man/> - e.g. <http://linux.die.net/man/2/fork>

In this class, you'll be expected to consult the Linux manual pages and to do some reading and research on your own, so practice this in this first lab and try to answer as many of your own questions as possible, but do come to office hours and ask for help if you get stuck.

Upload all code and your report completed using MS Word or as a PDF to Canvas and include all source code (ideally example output should be integrated into the report directly, but if not, clearly label in the report and by filename if test and example output is not pasted directly into the report). *Your code must include a Makefile so the TAs can build your solution on Ubuntu VB-Linux, a Jetson, a Raspberry Pi or a Altera DE1-SoC. Please zip or tar.gz your solution with your first and last name embedded in the directory name and/or provide a GitHub public or private repository link. Note that I may ask you or SA graders may ask you to walk-through and explain your code. Any code that you present as your own that is "re-used" and not cited with the original source is plagiarism. So, be sure to cite code you did not author and be sure you can explain it in good detail if you do re-use, you must provide a proper citation and prove that you understand the code you are using.*

Grading Rubric

1) [10 points] Priority inheritance paper review and unbounded priority inversion issues:

Section Evaluation	Points Possible	Score	Comments
Priority Inheritance Paper Summary with 3 key points	3		
Does the PI-futex provide safe and accurate protection from unbounded priority inversion? If not, what is different about it?	2		
Why did some of the Linux community resist addition of priority inversion solutions?	2		
Position on which solution is best for use with Linux	3		
Total	10		

2) [25 points] Thread safety with global data and issues for real-time services:

Section Evaluation	Points Possible	Score	Comments
Describe clearly what it means to write "thread safe" functions that are "re-entrant".	5		
Describe each method and impact to real-time threads/tasks.	10		
Correct shared state update code	10		
Total	25		

3) [20 points] Example synchronization code using POSIX threads:

Section Evaluation	Points Possible	Score	Comments
Demonstration and description of fixed deadlock with threads	5		
Demonstration and description of priority inversion with threads	5		
Description of RT_PREEMPT_PATCH and whether Linux can be made real-time safe	5		

Does it make sense to simply switch to an RTOS and not use Linux at all for both HRT and SRT services?	5		
Total	20		

4) [15 points] Example synchronization code using POSIX message queues and threads:

Section Evaluation	Points Possible	Score	Comments
Demonstration of message queues on target board adapted from examples	10		
Description of how message queues would or would not solve issues associated with global memory sharing	5		
Total	15		

5) [20 points] Watchdog timers (for the system) and timeouts (for API calls):

Section Evaluation	Points Possible	Score	Comments
Describe how Linux WD timer might be used if software caused an indefinite deadlock	10		
Adaptation of code from #2 MUTEX sharing to handle timeouts for shared state	10		
Total	20		

6) [10 points] TA Review

Section Evaluation	Points Possible	Score	Comments
Individually answer questions about the Exercise sections 1-5	5		
Demonstrate to the TAs that your code for section 2 through 5 is complete and working.	5		
TOTAL	10		