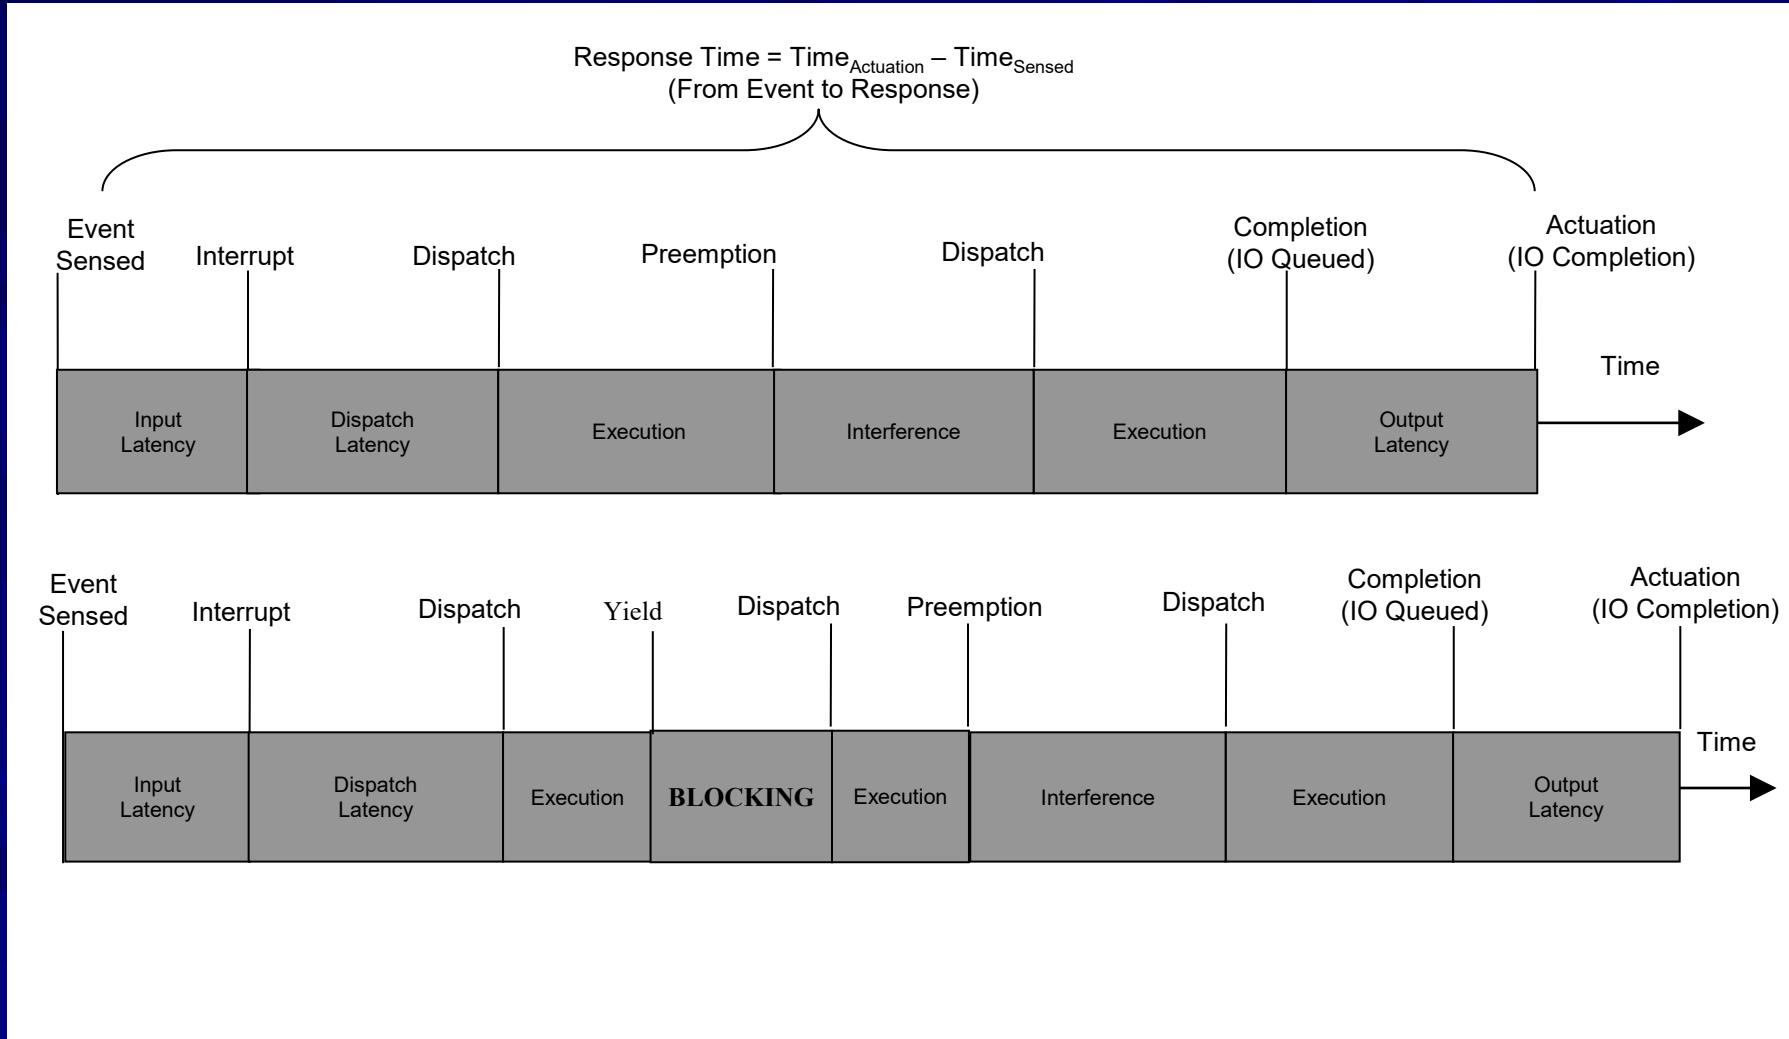


# **ECEN 5623**

## **IO Services**

**Be Boulder.**  
 University of Colorado **Boulder**

# Service Response Timeline (With Intermediate Blocking)



# Services and I/O

## ■ Initial Input I/O from Sensors

- Low-Rate Input: Byte or Word Input from Memory-Mapped Registers or I/O Ports
- High-Rate Input: Block Input from DMA Channels and Block Reads from FIFOs

## ■ Final Response I/O to Actuators

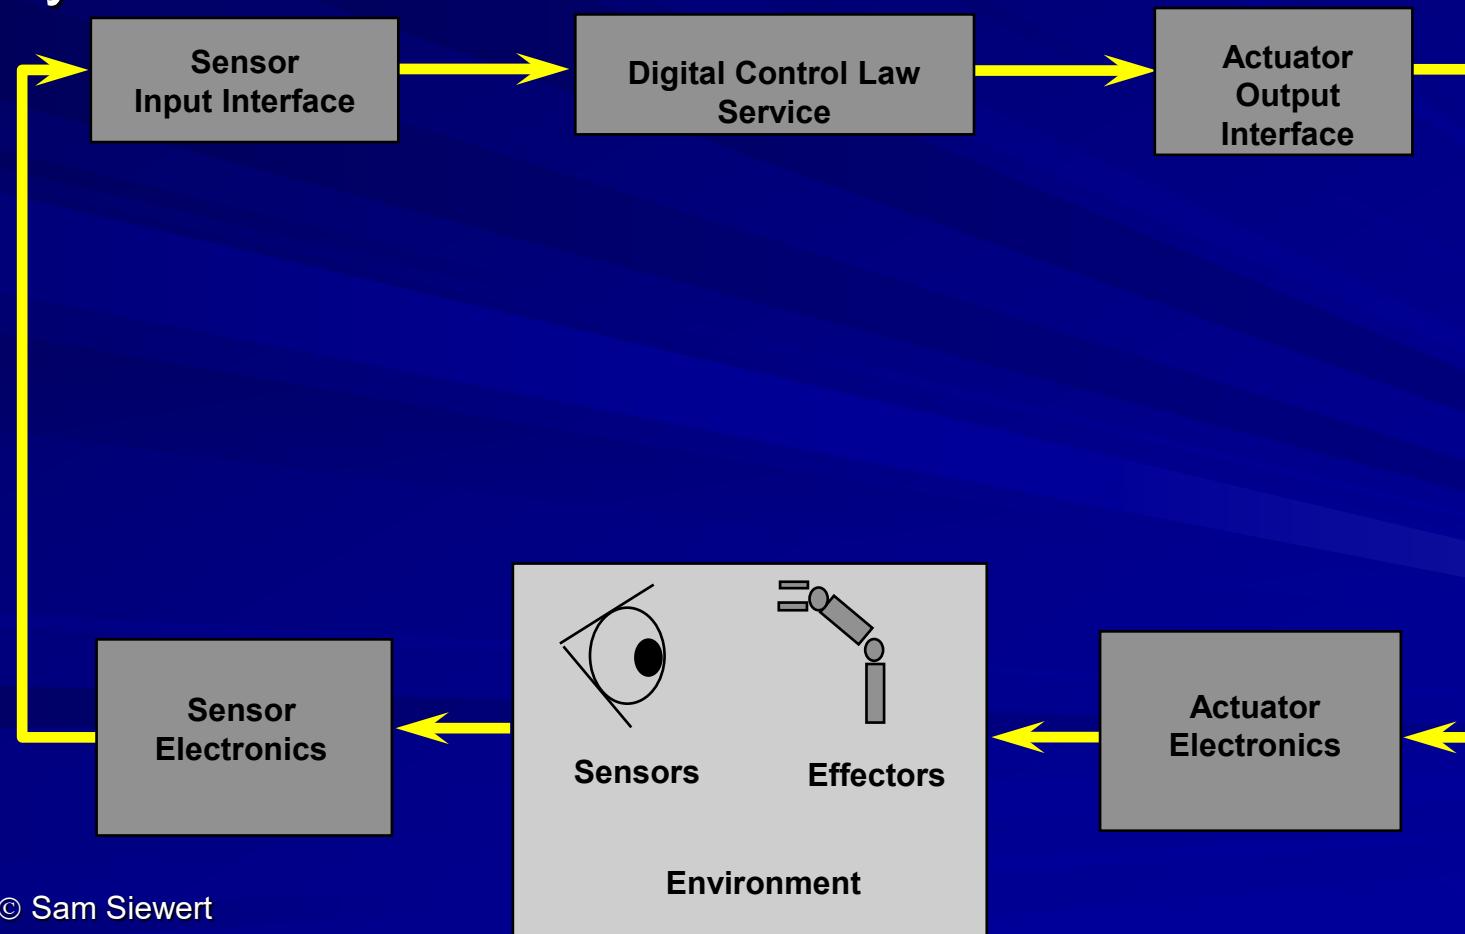
- Low-Rate Output: Byte or Word Writes Posted to Output FIFO or Bus Interface Queue
- High-Rate Output: Block Output on DMA Channels and Block Writes

## ■ Intermediate I/O

- During Service Execution, Memory Mapped Register I/O
- External Memory I/O
- Cache Loads and Write-Backs
- In Memory Input/Output from Service to Service

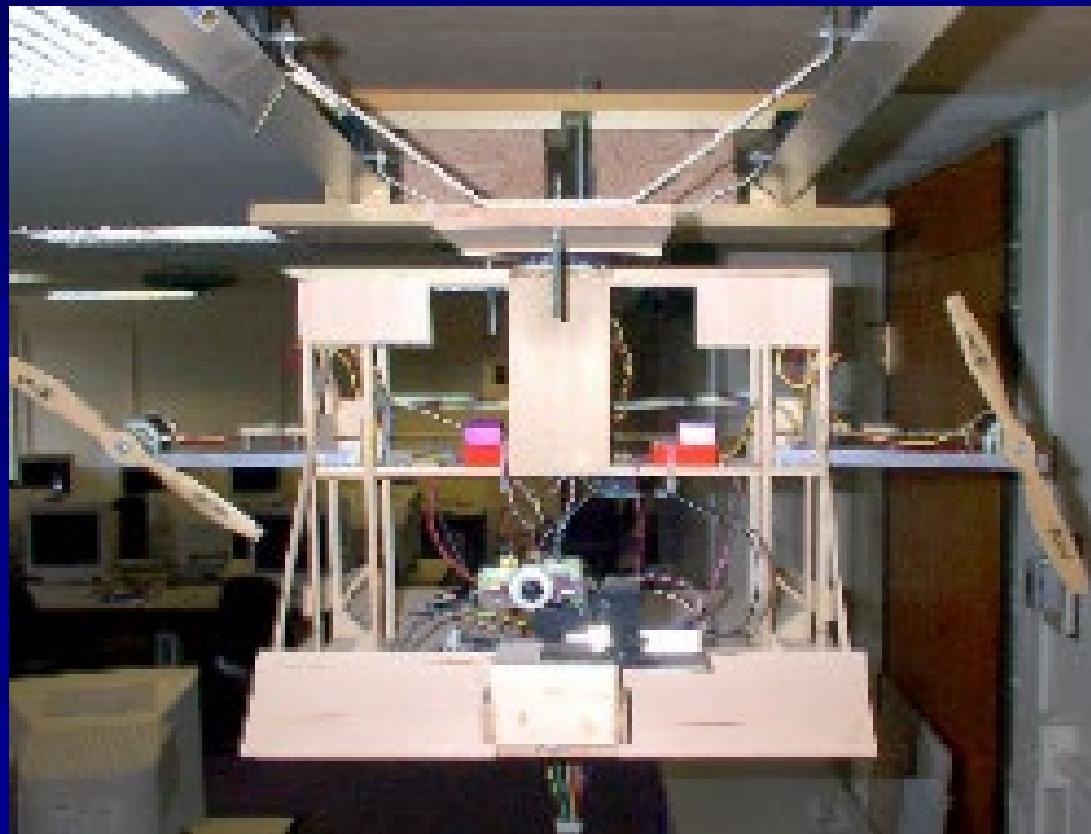
# Service Integration Concepts

- Initial Sensor Input and Final Actuator Response Output
  - Device Interface Drivers
- Simple Service Has No Intermediate I/O Latency, No Synchronization

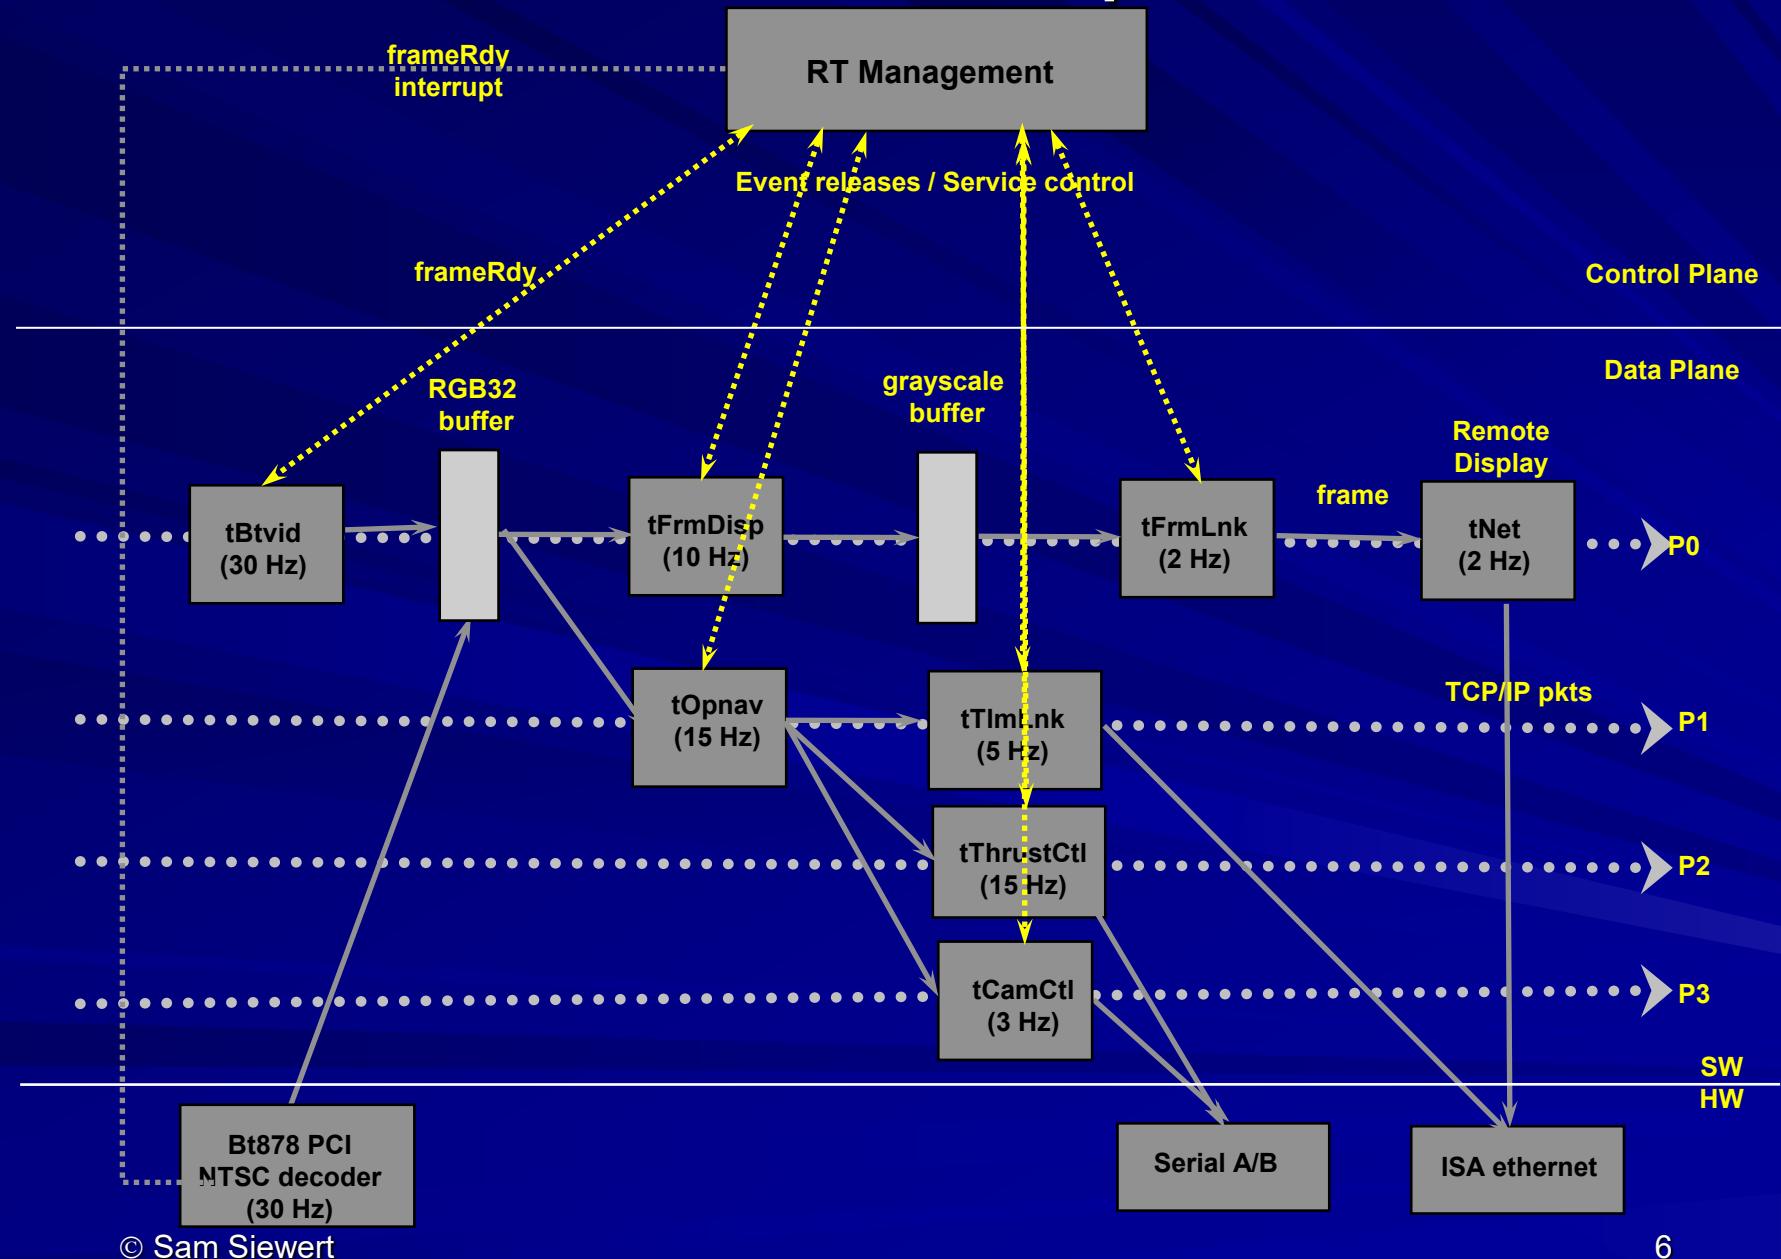


# Complex Multi-Service Systems

- Multiple Services
- Synchronization Between Services
- Communication Between Services
- Multiple Sensor Input and Actuator Output Interfaces
- Intermediate I/O, Shared Memory, Messaging

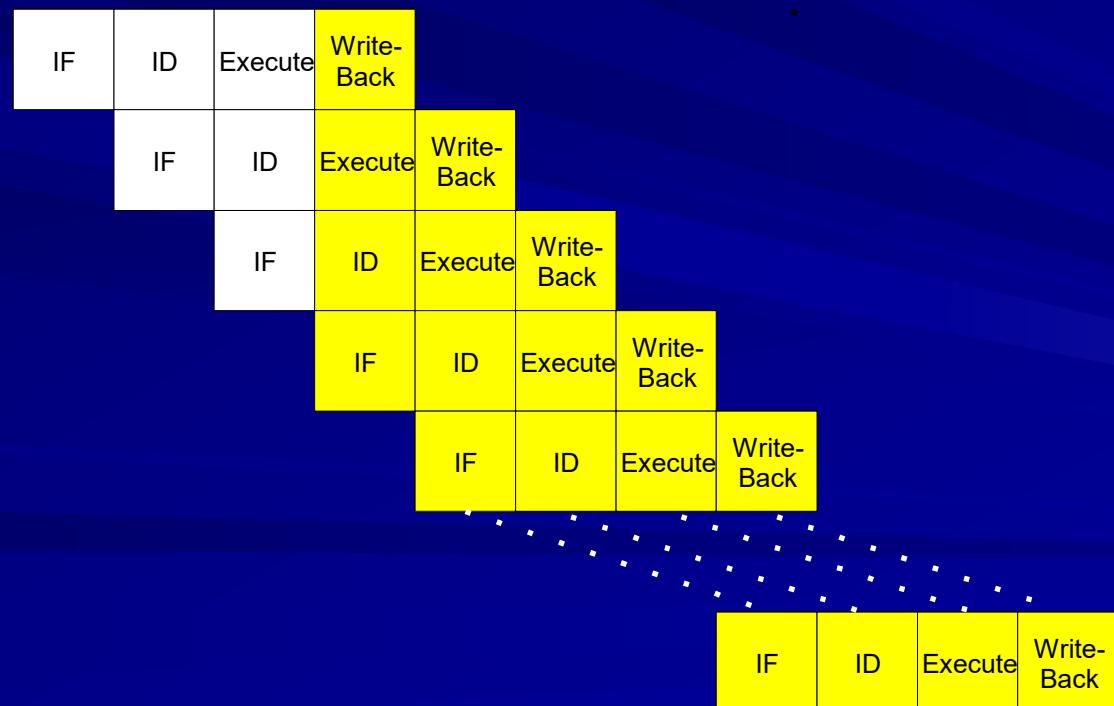


# Multi-Service Pipelines



# Pipelined Architecture Review

- Recall that Pipeline Yields CPI of 1 or Less
- Instruction Completed Each CPU Clock
- Unless Pipeline Stalls!



# Service Execution

$$WCET = [(CPI_{best-case} \times Longest\_Path\_Inst\_Count) + Stall\_Cycles] \times Clk\_Period$$

- Efficiency of Execution
- Memory Access for Inter-Service Communication
- Bounded Intermediate I/O
- Ideally Lock Data and Code into Cache or Use Tightly Coupled Memory [Scratch Pad, Zero Wait State]

# Service Efficiency

## ■ Path Length for a Release

- Instruction Count
- Longest Path Given Algorithm
  - Branches
  - Loop Iterations
  - Data Driven?

## ■ Path Execution

- Number of Cycles to Complete Path
- Clocks Per Instruction
- Number of Stall Cycles for Pipeline
  - Data Dependencies (Intermediate IO)
  - Cache Misses (Intermediate Memory Access)
  - Branch Mis-predictions (Small Penalty)

# Hiding Intermediate IO Latency (Overlapping CPU and Bus I/O)

- ICT = Instruction Count Time
  - Time to execute a block of instructions with no stalls
  - CPU Cycles x CPU Clock Period
- IOT = Bus Interface IO Time
  - Bus IO Cycles x Bus Clock Period
- OR = Overlap Required
  - Percentage of CPU cycles that must be concurrent with I/O cycles
- NOA = Non-Overlap Allowable for  $S_i$  to meet  $D_i$ 
  - Percentage of CPU cycles that can be in addition to IO cycle time without missing service deadline
- $D_i$  = Deadline for Service  $S_i$  relative to release
  - interrupt or system call initiates  $S_i$  request and  $S_i$  execution
- CPI = Clocks Per Instruction for a block of instructions
  - IPC is Instructions Per Clock, Also Used, Just the Inverse (Superscalar, Pipelined)

# Processing and I/O Time Overlap

- $D_i \geq IOT$  is required; otherwise if  $D_i < IOT$ ,  $S_i$  is ***IO-Bound***
- $D_i \geq ICT$  is required; otherwise if  $D_i < ICT$ ,  $S_i$  is ***CPU-Bound***
- $D_i \geq (IOT + ICT)$  **requires no overlap** of IOT with ICT
- if  $D_i < (IOT + ICT)$  where ( $D_i \geq IOT$  and  $D_i \geq ICT$ ), **overlap of IOT with ICT is required**
- if  $D_i < (IOT + ICT)$  where ( $D_i < IOT$  or  $D_i < ICT$ ), deadline ***Di can't be met regardless of overlap***

# Service Execution Efficiency – Waiting on I/O from MMIO Bus and Memory

- $CPI_{worst-case} = (ICT + IOT) / ICT$

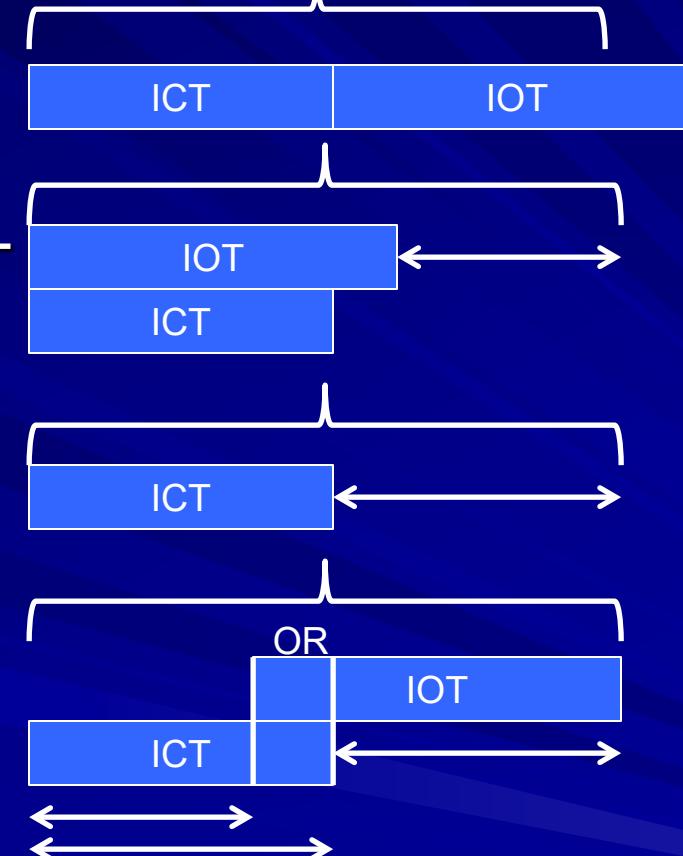
- $CPI_{best-case} = (\max(ICT, IOT)) / ICT$

- $CPI_{required} = D_i / ICT$

- $OR = 1 - [(D_i - IOT) / ICT]$

- $CPI_{required} = [ICT(1-OR) + IOT] / ICT$

- $NOA = (D_i - IOT) / ICT; OR + NOA = 1 \text{ (by definition)}$



# Synchronization and Message Passing

## ■ Message Queues

- Provide Communication
- Provide Synchronization

## ■ Traditional Message Queue

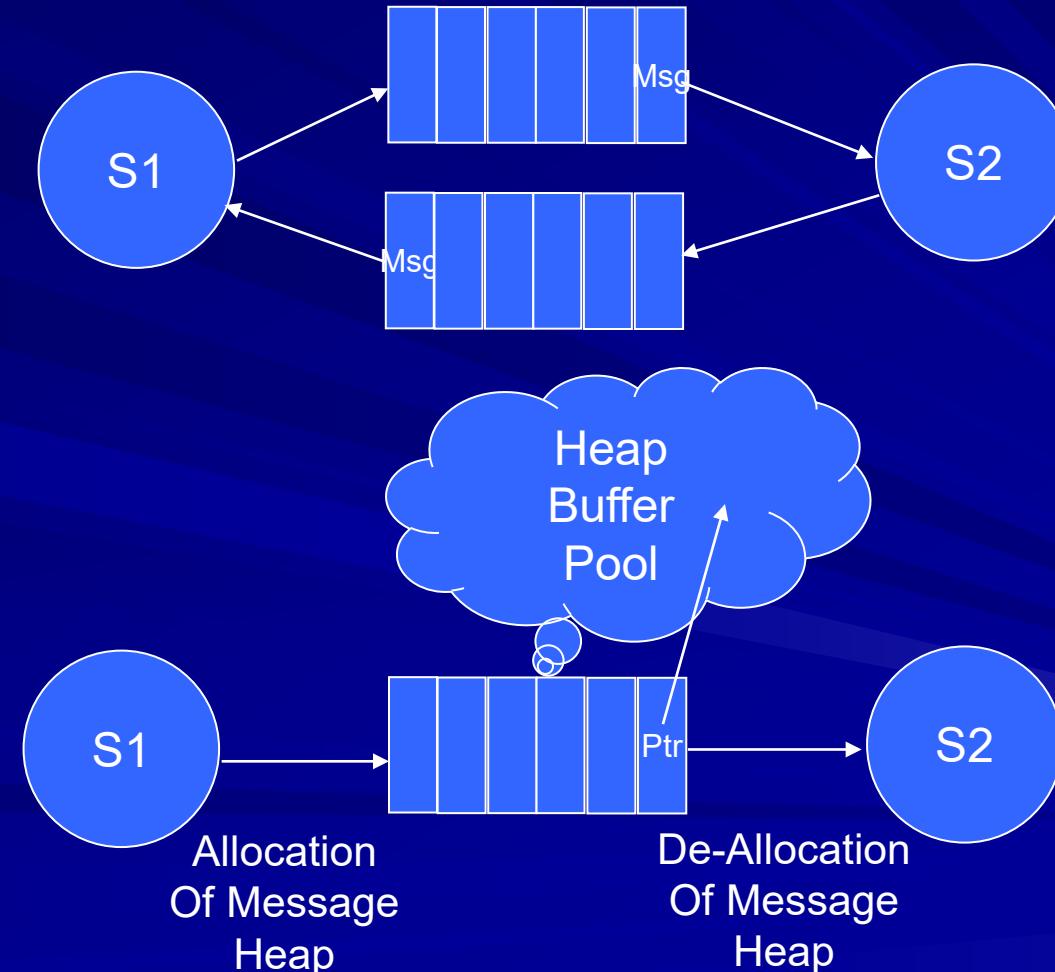
- Message Size
- Internal Fragmentation

## ■ Heap Message Queue

- Messages Are Pointers
- Message Data in Heap

### ENEA OSE

- Message Passing RTOS
- Advertised Advantages



# Synchronization and Message Passing

## ■ Message Queues

- Provide Communication
- Provide Synchronization

## ■ Traditional Message Queue

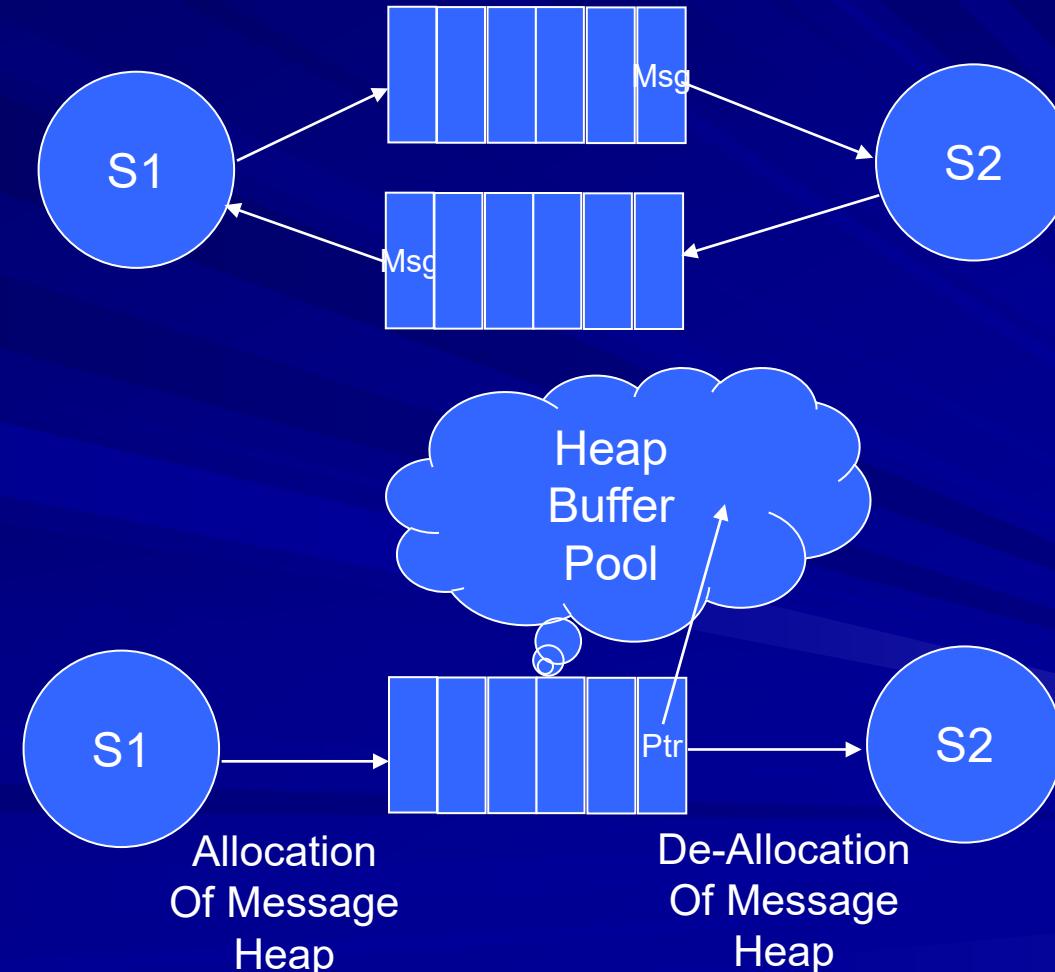
- Message Size
- Internal Fragmentation

## ■ Heap Message Queue

- Messages Are Pointers
- Message Data in Heap

### ENEA OSE

- Message Passing RTOS
- Advertised Advantages



# What About Storage I/O

- Much Higher Latency than Off-Chip Memory or Bus MMIO
- Milliseconds of Latency Compared to GHz Core Clock Rates – 1 million to 1!
- Flash Better, But Still Slower than SRAM or DRAM
- Option #1 – Avoid Disk Drives and Flash I/O
- Option #2 – Pre-Buffer Storage I/O (Read-Ahead Cache), Post Write-backs (Write-Back Cache), MFU/MRU (Most Frequently Used / Most Recently Used) Kept in Read Cache

# Parallel Processing Speed-up (Makes I/O Latency Worse?)

- Grid Data Processing Speed-up
  1. Multi-Core, Multi-threaded, Macro-blocks/Frames
  2. SIMD, Vector Instructions Operating over Large Words (Many Times Instruction Set Size)
  3. Co-Processor Operates in Parallel to CPU(s)

- SPMD – GPU or GP-GPU Co-Processor
  - PCI-Express Bus Interfaces
  - Transfer Program and Data to Co-Processor
  - Threads and Blocks to Transform Data Concurrently

- Image Data Processing – Few Data Dependencies
  - Good Speed-up by Amdahl's Law
  - P=Parallel Portion
  - (1-P)=Sequential Portion
  - S=# of Cores (Concurrency)
  - Overhead for Co-Processor
  - IO for Co-Processing

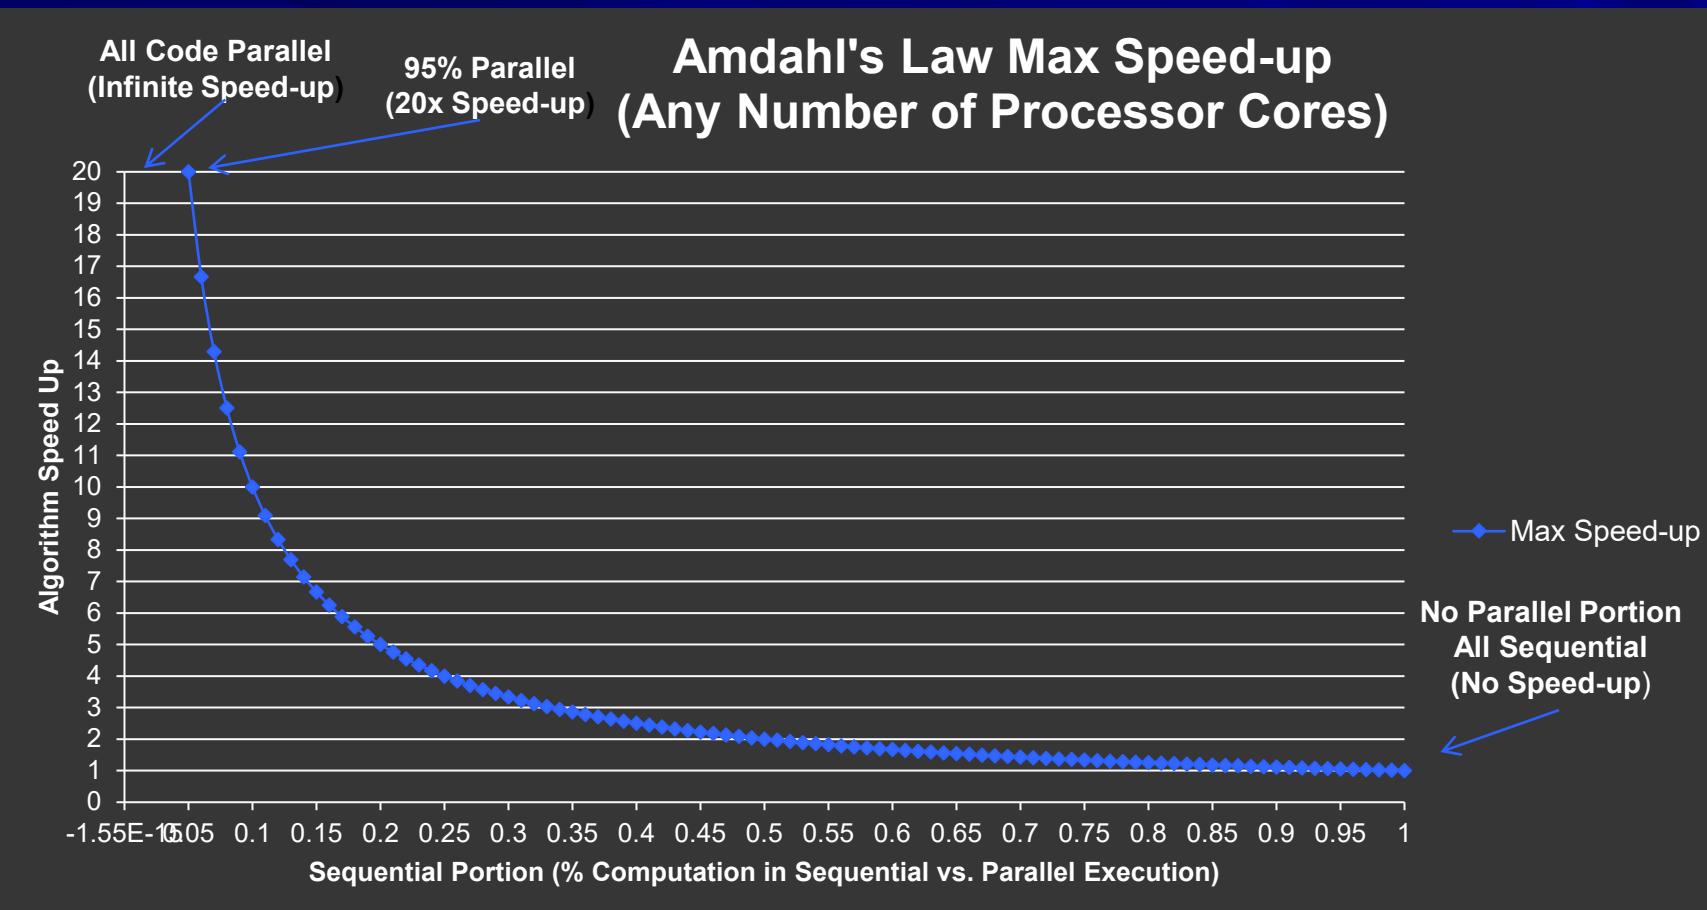
S is infinite here

$$\text{Max\_Speed\_Up} = \frac{1}{(1-P)+0}$$

$$\text{Multicore\_Speed\_Up} = \frac{1}{(1-P)+P/S}$$

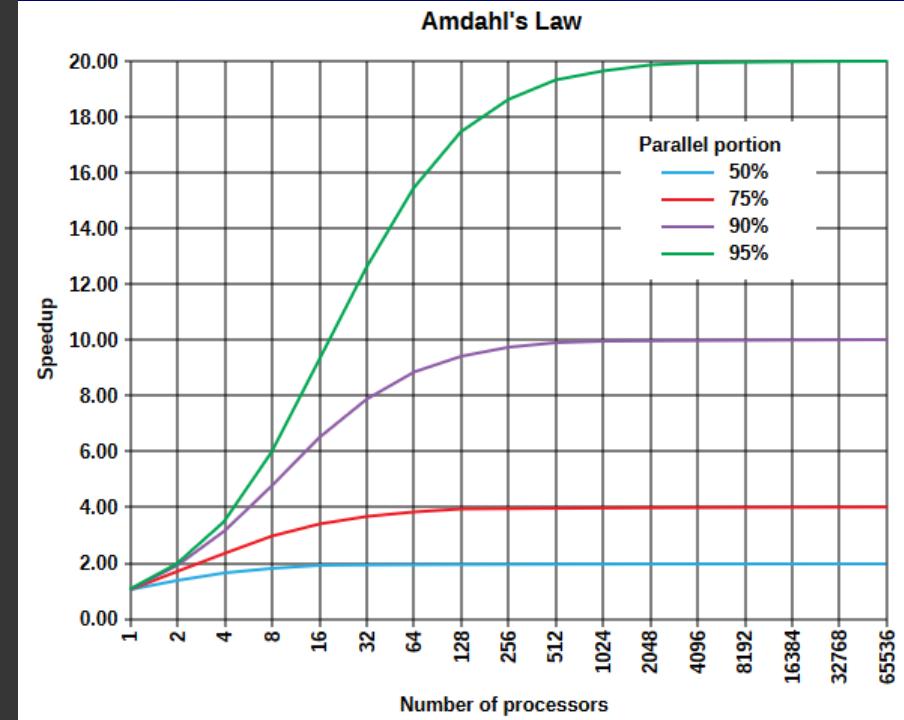
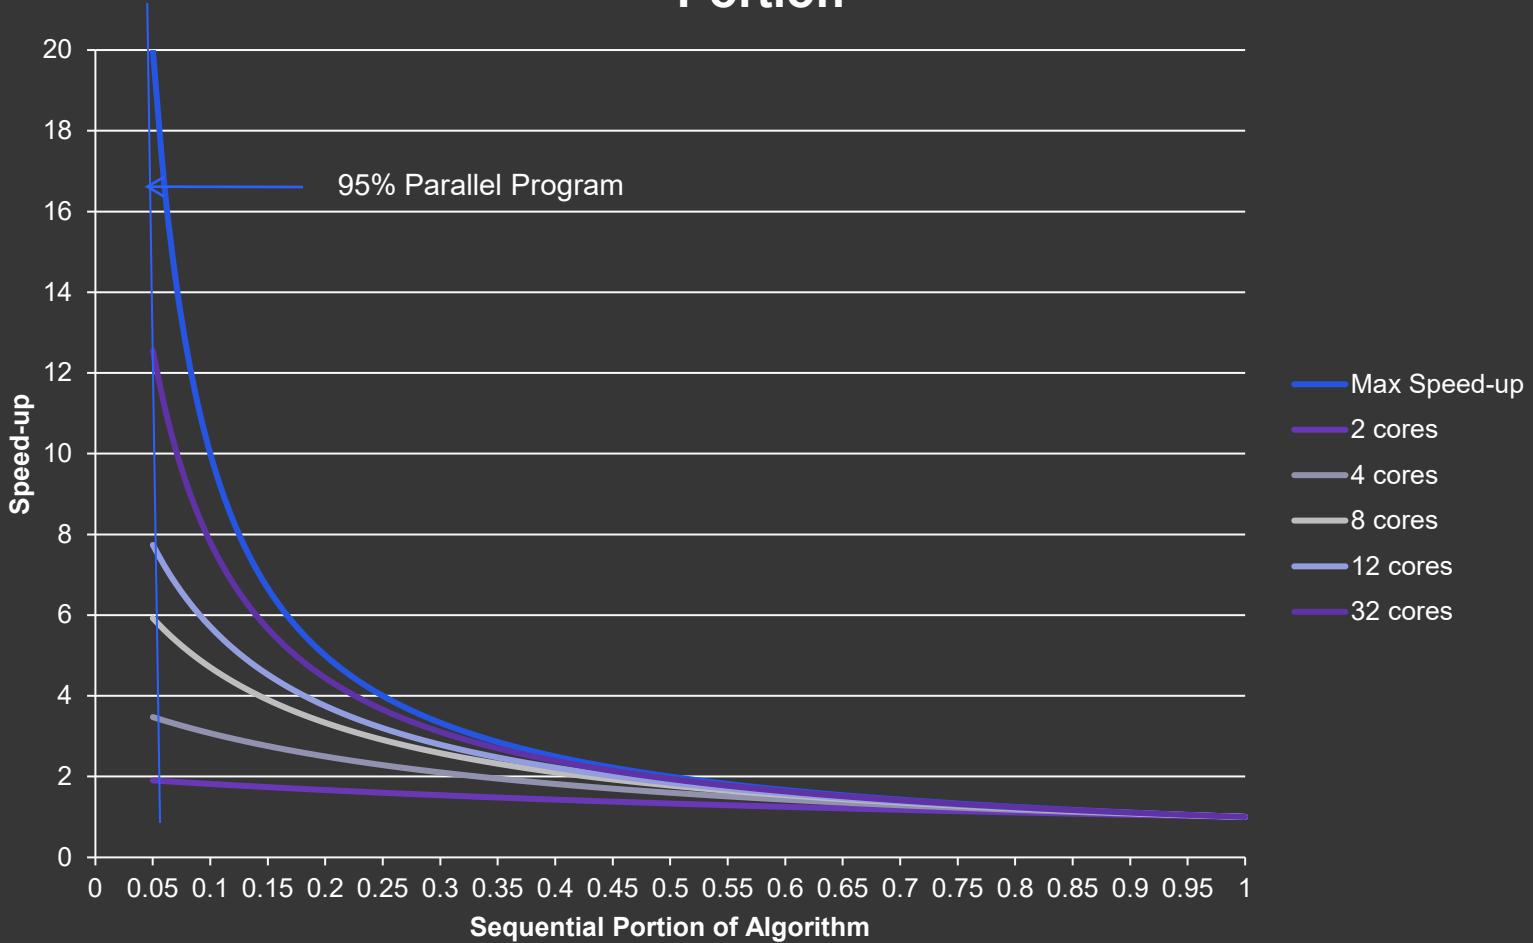
# Amdahl's Law – Infinite Cores

- Maximum Speed-up Driven by Sequential and Parallel Portions of Program
  - $P$  = Parallel Portion
  - $(1-P)$  = Sequential Portion
  - Speed-up for Given Multi-core Architecture Function of # of Cores (Speed-up in Parallel Portions)



# Multi-Core Speed-Up

Amdahl's Law - Speed-up with # Cores and Parallel Portion



# Hiding Storage I/O Latency – Overlapping with Processing

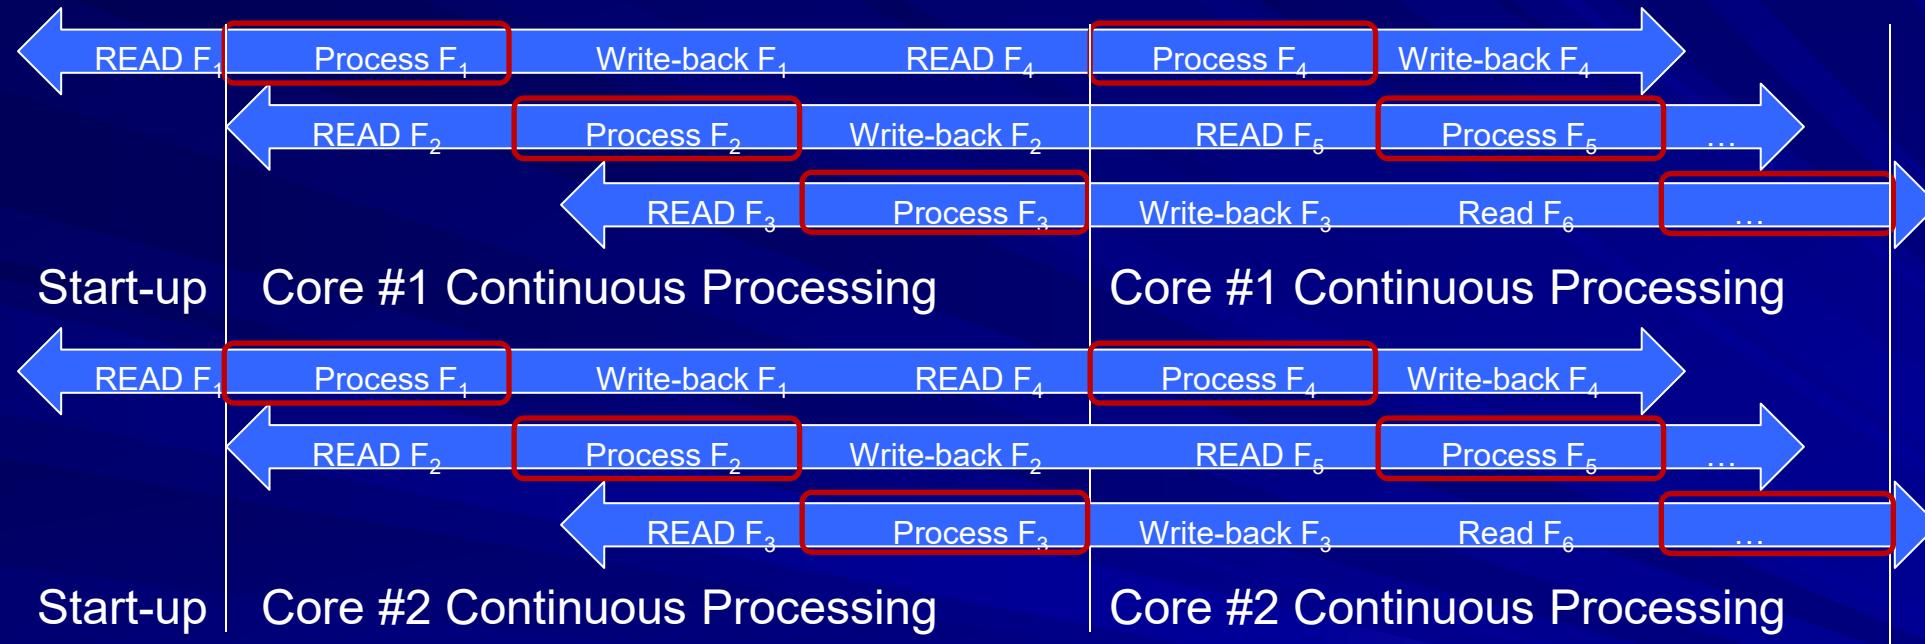
- Simple Design – Each Thread has READ, PROCESS, WRITE-BACK Execution



- Frame rate is READ+PROCESS+WRITE latency – e.g. 10 fps for 100 milliseconds
  - If READ is 70 msec, PROCESS is 10 msec, and WRITE-BACK 20 msec, predominate time is IO time, not processing
  - Disk drive with 100 MB/sec READ rate can only read 16 fps, 62.5 msec READ latency

# Hiding Storage I/O Latency

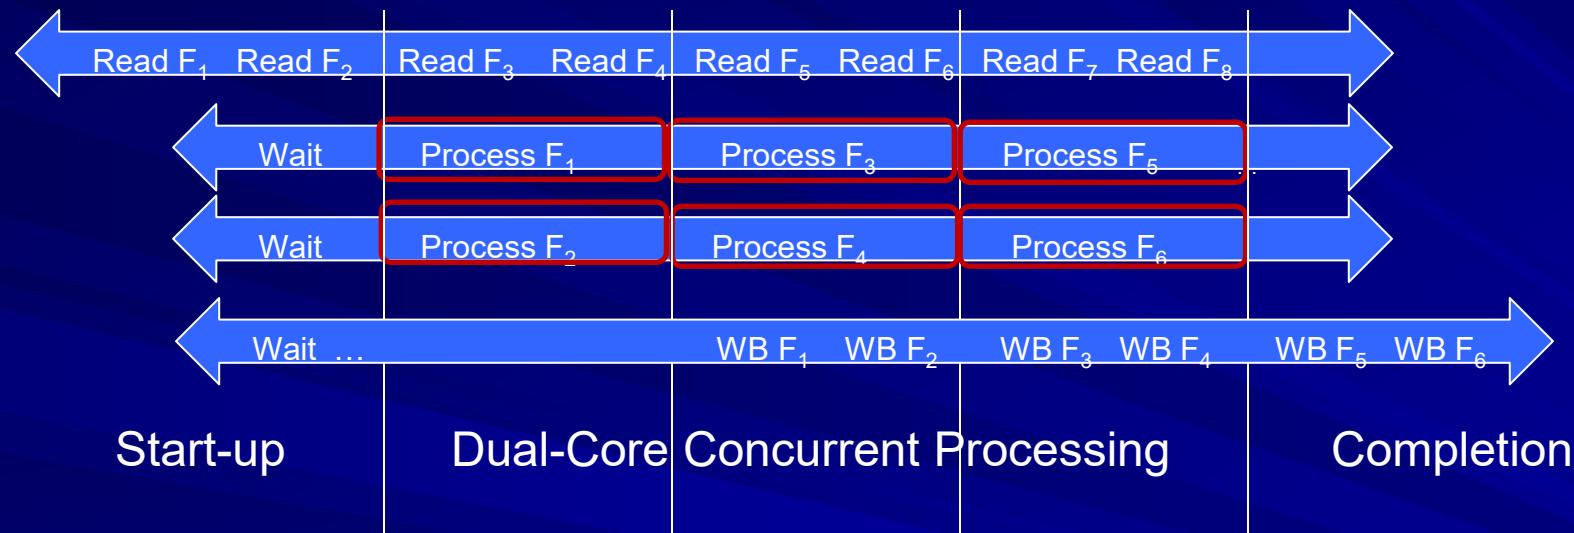
## ■ Schedule Multiple Overlapping Threads?



- Requires  $N_{\text{threads}} = N_{\text{stages}} \times N_{\text{cores}}$
- 1.5 to 2x Number of Threads for SMT (Hyper-threading)
- For IO Stage Duration Similar to Processing Time
- More Threads if IO Time (Read+WB+Read)  $\gg 3 \times$  Processing Time

# Hiding Latency – Dedicated I/O

## ■ Schedule Reads Ahead of Processing



- Requires  $N_{\text{threads}} = 2 + N_{\text{cores}}$
- Synchronize Frame Ready/Write-backs
- Balance Stage Read/Write-Back Latency to Processing
- 1.5 to 2x Threads for SMT (Hyper-threading)

# Processing Latency Alone

## ■ Write Code with Memory Resident Frames

- Load Frames in Advance
- Process In-Memory Frames Over and Over
- Do No IO During Processing
- Provides Baseline Measurement of Processing Latency per Frame Alone
- Provides Method of Optimizing Processing Without IO Latency

# I/O Latency Alone

- Comment Out Frame Transformation Code or Call Stubbed NULL Function
  - Provides Measurement of IO Frame Rate Alone
  - Essentially Zero Latency Transform
  - No Change Between Input Frames and Output Frames
  - Allows for Tuning of IO Scheduler and Threading

# Tips for Linux I/O Scheduling

- `blockdev --getra /dev/sda`
  - Should return 256
  - Means that reads read-ahead up to 128K
  - Function calls – read, fread should request as much as possible
  - Check “actual bytes read”, re-read as needed in a loop
- `blockdev --setra /dev/sda 16384 (8MB)`
- **Switch CFQ to Deadline**
  - Use “lsscsi” to verify your disk is /dev/sda ... substitute block driver interface used for file system if not sda
  - `cat /sys/block/sda/queue/scheduler`
  - `echo deadline > /sys/block/sda/queue/scheduler`
- Options are noop, cfq, deadline, anticipatory